

[Open in app](#)**Jonathan**[Follow](#)

29 Followers

[About](#)

Certified Kubernetes Security Specialist (CKS) Preparation Part 4— Cluster Hardening



Jonathan Feb 24 · 9 min read

If you have not yet checked the previous parts of this series, please go ahead and check [Part1](#), [Part2](#) and [Part3](#).

In this article, I would focus on the preparation around **cluster hardening** in CKS certification exam.

Role and Role Binding

- Role = the position that could perform actions
- RoleBinding = the binding of user/service account and roles
- Roles are namespace specific.

Let's see an example to get a clearer idea. Our goal is to create a role that could get pods in namespace "test" and bind user "jon" to this role.

Create a Role called "get-pod" in namespace "test"

- `kubectrl create role get-pod — verb=get — resource=pods -n test`

```
jonw@CKS-Master:~$ kubectl create role get-pod --verb=get --resource=pods -n test
I0203 23:22:33.676586 11005 request.go:655] Throttling request took 1.143225893s, request: GET:https://192.168.1.4:6443/apis/coordination.k8s.io/v1?timeout=32s
role.rbac.authorization.k8s.io/get-pod created
```

[Open in app](#)

- *kubectl create rolebinding get-pod-jon — role=get-pod — user=jon -n test*

```
jonw@CKS-Master:~$ kubectl create rolebinding get-pod-jon --role=get-pod --user=jon -n test
rolebinding.rbac.authorization.k8s.io/get-pod-jon created
```

Test whether user “jon” could actually get pods in namespace “test”

- *kubectl auth can-i get pods — as jon -n test*
- *kubectl auth can-i get pods — as jon -n default*

```
jonw@CKS-Master:~$ kubectl auth can-i get pods --as jon -n test
yes
jonw@CKS-Master:~$ kubectl auth can-i get pods --as jon -n default
no
```

Cluster Role and Cluster Role Binding

- ClusterRole = the position that could perform actions across the whole cluster
- ClusterRoleBinding = the binding of user/service account and cluster roles
- ClusterRoles are NOT namespace specific.

Let’s see an example to have a clearer idea. Our goal is to create a cluster role that could delete pods in namespace “test2” and bind user “jon” to this cluster role.

Create a ClusterRole called “delete-pod”

- *kubectl create clusterrole delete-pod — verb=delete — resource pod*

```
jonw@CKS-Master:~$ kubectl create clusterrole delete-pod --verb=delete --resource pod
I0203 23:28:19.769156 14677 request.go:655] Throttling request took 1.109179257s, request: GET:https://192.168.1.4:6443/apis/rbac.authorization.k8s.io/v1beta1?timeout=32s
clusterrole.rbac.authorization.k8s.io/delete-pod created
```

Create a ClusterRoleBinding called “delete-pod-jon” that connects the ClusterRole “delete-pod” with user “jon”

- *kubectl create clusterrolebinding delete-pod-jon — clusterrole delete-pod — user jon*

[Open in app](#)

Test whether user “jon” could actually delete pods in all namespaces.

- *kubectl auth can-i delete pods — as jon -n test*
- *kubectl auth can-i delete pods — as jon -n default*

```
jonw@CKS-Master:~$ kubectl auth can-i delete pods --as jon -n test
yes
jonw@CKS-Master:~$ kubectl auth can-i delete pods --as jon -n default
yes
```

Certificate Signing Requests

Certificate signing requests (CSR) are essentially users or service accounts asking kube-apiserver to provide access for managing K8s clusters with their own identity. This would ensure that users and service accounts would be communicating with kube-apiserver on their own behalf when having interactions. The basic process flow is

User/service account generates a key

- *openssl genrsa -out jon.key 2048*

```
jonw@CKS-Master:~$ openssl genrsa -out jon.key 2048
Generating RSA private key, 2048 bit long modulus (2 primes)
.+++++
.....+++++
e is 65537 (0x010001)
```

Use the key to generate a CSR

- *openssl req -new -key jon.key -out jon.csr*

```
jonw@CKS-Master:~$ openssl req -new -key jon.key -out jon.csr
Can't load /home/jonw/.rnd into RNG
140466202685888:error:2406F079:random number generator:RAND_load_file:Cannot open file:../crypto/rand/
randfile.c:88:Filename=/home/jonw/.rnd
You are about to be asked to enter information that will be incorporated
into your certificate request.
What you are about to enter is what is called a Distinguished Name or a DN.
There are quite a few fields but you can leave some blank
For some fields there will be a default value,
If you enter '.', the field will be left blank.
-----
Country Name (2 letter code) [AU]:
State or Province Name (full name) [Some-State]:
Locality Name (eg. city) []:
```

[Open in app](#)

```
Please enter the following 'extra' attributes
to be sent with your certificate request
A challenge password []:
An optional company name []:
```

Encode the output CSR with base 64 and copy the content to a K8s CSR YAML. Check [here](#) for getting the default K8s CSR YAML template.

- nano k8s-jonw-csr.yaml
- cat jon.csr | base64 -w 0

```
apiVersion: certificates.k8s.io/v1
kind: CertificateSigningRequest
metadata:
  name: jon
spec:
  groups:
  - system:authenticated
  request: LS0tLS1CRUdJTiBDRVJUSUZJQ0FURSBRSRVFVRVNULS0tLS0KTULJQ21EQ0NBWUFQVFBd1V6RUxNQWtHQTFVRUJJoTUN
RVLV4RXpBUBk5bnJlZCQWdWQ2X0dmJXVXRVM1JoZEdVeApJVEFmQmdOVkKBb01HRWx1ZEdweWJtVjBJRmRwWkdKcGRITWdVSFI1SUUV4M
FpERU1NQW9HQTFVRUF3d0RhbTl1ck1JSUJJakFOQmdrcwRaUc5d2BCQVFFRkFBT0NBUThtTUJQKNNs0NB0UUVBMi81QjBNU1prY2V
4K1NaSElsQWMMKNE9FTGRtQXA2K2owekY4cXZkYkY5VVBnSk1WbkE5eGczNFdGb1VlWmpqaEY2eTZubDlySFByRHhLOWhaRGRSbApTT
UdDN09Q0itzaGNJMNuU40EF4YXVDMlowVXJCbWZRMklnUHQVd40Eg0R0U1UvG0c1YwdmLR0G11ZXBjTm44ClBoRk3peFhPYm9xUU00
vZXJaTGpranIwTXFBU1F5M1BvNjhPQmhhVGJZYk5KeKpya1BTTWpncTM4d1F2YH9EMmgK0HA2N3dxVXlINHYxUXNVZxpLZ1N0TndRM
k5van50T1lhVmdJRzNPMEhdChRoYWo2dUU3ZG9BTkZzXpV3RUUWpENHV3eHR2S003RmorZEXjUGZ5Q3R2UEQyWitCSmF6aGRxNjk
0THFDM1VQ0k52UWZvck9GZjh3VThncTM0U3F4CnR3SURBUUFcb0FBd0RRWUpLb1pJaHZjTkFRRUxCUUFEZ2dFQkFKS0xSRWlMaER6V
EFFTmR5NEtPUUVEOHJjT2YKRVBKnlFvQXdGYUUV0QUV3MzJRvYy9YbE9iYUw1ZWJ2a01ITLR5KzcZnURVQ2tms0VRTGJQZzJyZ05tN3V
hclhjeQpsZy9CakovcGRIQlIza2Fia050czQ0TUZEQmFpaldPZ2s4dVZQeGpNNzVsY1VcdVl3aC9icWpxcys2RGlVT0k3Cm9zSlRku
0xnakdBWtJYc1hLUG9tcE9CV1Rkak1oaFZ6aStSbG0xdEpDUeViW0Nac3JXWnV2ZU1hekZtQVFaaE4KYnV1cngxRlRtNTk0YkhbG1
wbGZoT0V4MTBEOGxNUzVnYmlywEteA3JoNVR0azJWbU1wcE1WRzZmaGtEh1diMqpvTTRLUW52YzZsUFFmaH2MRGtVL3NUTzNRcG9ZN
DRCZmRLTz16SmJUaHFLbik4ZUFvVNVNmdyV2xFT0KL50tLS1FTkQgQ0VSVElGSUNBVEUgUkVRVUUVTC0tLS0tCg==
  signerName: kubernetes.io/kube-apiserver-client
  usages:
  - client auth
```

Create the K8s CSR

- `kubectl create -f k8s-jon-csr.yaml`

```
jonw@CKS-Master:~$ kubectl create -f k8s-jon-csr.yaml
certificatesigningrequest.certificates.k8s.io/john created
```

Check CSR status and Approve the CSR

- `kubectl get csr`
- `kubectl certificate approve jon`

```
jonw@CKS-Master:~$ kubectl get csr
```

NAME	AGE	SIGNERNAME	REQUESTOR	CONDITION
------	-----	------------	-----------	-----------

[Open in app](#)


NAME	AGE	SIGNERNAME	REQUESTOR	CONDITION
jon	27s	kubernetes.io/kube-apiserver-client	kubernetes-admin	Approved, Issued

Get client certificate from K8s CSR YAML and Decode it from base 64 and save it to a new file

- `kubectl get csr jon -o yaml`
- `echo <certificate content> | base64 -d > jon.crt`
- `cat jon.crt`

```
jonw@CKS-Master:~$ echo "LS0tLS1CRUdJTiBDRVJUSUZZJQ0FURSB0tLS0tCk1JSURPVENDQWlHZ0F3SUJBZ0lScUw1Mw5MZWcyQW5CYmhoQmdCdGNpcXN3RFFZSkVWklodmNOQVFFTEJRQXcKRLRFVE1CRUdBMVVQXhNS2EzVmLaWEp1WlhSbGN6QWVgdzB5TVRBeU1EUXdNREU0TWp0YUZZ3MHlNakF5TURRdWpNREU0TWp0YU1GTxhDekFKQmdOVKBWVRBa0ZWTWJnd0VRURWUUVJRXxwVGIyMwVxMwVWk4WVhSbE1TRXIdId1lECLZRUUtFeGhKYm5SbGNTNwXkQ0JYYVdSbmFYUnpJRkIwZVNCVWRHUXhEREFLQmdOVKBWVRBmNBM2YmpDQ0FTSxcKRFFZSkVWklodmNOQVFFQkJRQURnZ0VQURDQ0FRb0NnZ0VCQU52K1FkREVtWkhIc2ZrbVJ5S1FIT0RoQzNaZwpLZXZvOU14ZktyM1d4Z1dqNENTRlp3UGNZT1tGaFoxQjJZNDRSZXN1cDVMYXh6Nnc4U3ZVZVIEZVVpVakJndXpqCndmcklyQ05yd1BBTVdyZ3RtZEZld1pueME5pRytSNHdGc2ZCK0JBbmtGK0xGZEw0a1BKcm5xU0RaL0Q0U1Fzc1YKcm02S2tEUDNkMlM0NUk2OUURLZ0VrRWRqMU92RgdZV2syMkd6U2N5YTR6MGpJMzZ0L0wwTDJ6ZzlvZktldThLbApNaCtMOVVMRkhzM29FaLRjRU5qYUk1eLRtR2xZQ0J0enRCd3FiwVdvK3JoTzNhQURRck5jMUJyVtBnK0xzTjWJiCnladXhZL25TM0QzOGdyYnp3OW1mZlNXczRYVXV2ZUM2Z3RsRHdUYjBINKt6aFgvTUZQSut0K0Vxc2JjQ0F3RUEKQWFOR01FUXdFd1lEVLlwbEJBd3dDZ1lJS3dZQkJRUVhBd0L3REFZRFZSMFRBUUgVqKfJd0FEQWZCZ05W5FNNRQpHREFXZ0JUR21iVUlnNHlHeC9MM1JTUnNjU3lRaC9DZDUQU5CZ2tXaGtpRz13MEJBUXNGQUFQ0FRRUFSdzg0ClpTVDRCN3hlcFNONXUyZ0k2b2JhbnN0HkxchhCLZQwblUwK3RWRDVs5ek1acl0mTdGYmlaSDJFT256M1pobzIKNTlkamM0UXNkckU4TWV3MUhod0FvdkV1NUVQMwtDdnJuZEXINXBQMS9Td0w5Mxo1QnVpZkVUULGwWxdGJTRApmlNSK3BiYVnPVXc3VGU0cXhibXNGTCS5bE1UYnVBcENxazR1YzRnYnVjUGJ0ckxntUV6dXNlGZ3di9XWkxHCmlpSTBlN1FLd0NBbFlJYjBzMLpqV2x0Y05QSFRkUm0xY2lXUVdybk9UM2RqWfJEQm9hc0lE2Nh0OUdvL2wYzckKaGozN2p5eHkyVmxzc210aG5GRj1PR0FUnzhCc1duRC9tdnAvTGNKL0NKNysveUVZQ09PamVJL3ltZkcrQVB5KwplKz12b0LVNktoRXLjQWZKOGc9PQ0tLS0tLUVORCBDRVJUSUZZJQ0FURSB0tLS0tCg==" | base64 -d > jon.crt
jonw@CKS-Master:~$ cat jon.crt
-----BEGIN CERTIFICATE-----
MIIDOTCCA1GgAwIBAgIRAL51nLfg2AnBbhhBgBtc1qwwDQYJKoZIhvcNAQELBQAw
FTETMBEGA1UEAxMKa3V1ZXJuZXRLczAeFw0yMTA5MDQwMDE4MjNaFw0yMTA5MDQw
MDE4MjNaMFwxCzAJBgNVBAYTAkFVMRMwEQYDVQIEwPTb21lLVN0YXRlMSEwHwYD
VQQKEWhhbnRlcmlkZCBXaWRnaXRzIFB0eSBMDGQxODAKBgNVBAMTA2pvcjCCASIw
DQYJKoZIhvcNAQEBBQADggEPADCCAQoCggEBANv+QdDEmZHHsfkmRyJQHODhC3Zg
Kevo9MxfK3WxfWj4CSFZwPcYN+FhZ1B2Y44Resup5faxz6w8SvYWQ3UZUjBguzj
wfrIXCNrvPAMWrgtmdFKWZn0NiG+R4wFsFB+BANKF+LFDL4kPjrnqSDZ/D4RQYsV
zm6KkDP3q2545I69DKgEkEdj10vDgYwk22GzScya4z0jI36t/L0L2zg90fKeu8Kl
Mh+L9ULFHs30EjTcEnjaI5zTmGLYCBtztBwqbyWo+rh03aADQrNc18rU0g+Lsmbb
yZuxY/nS3D38grbzW9mFSWs4XauveC6gtLDnTb0H6KzhX/MFPiKt+EqsbCawEA
AaNGMEQWewYDVR0LBAAwCgYIKwYBBQUHAWIwDAYDVR0TAQH/BAIwADAFBgNVHSME
GDAWgBTGmbUIM4yGx/L7RSR5ISyQh/CL6TANBgkqhkiG9w0BAQsFAAOCAQEARw84
ZST4B7XepSN5u2gI6n8gJpv8y1rxB/d0nU0+tpD5yzMZrWh17FbiZH2E0nz3Zho2
59dj4QsJrE8Mew1HhwAUvEu5EP1kCvrrndLH5pP1/SwL91z5BuiFEEQIFyhqtbsD
frSl+pbaSiUw7Te4qxbsmFL+9lNTbuApCqk4uc4gbcPbNrlGMEzusbxfwv/WZLG
iiI0e7QKwCALYIb0s2ZjwLtcNPHTdRm1ciwQWrrnOT3djXRDBoasIDfxt9Go/L/c7
hj37jxy2V1lsmthnFF9OGAT78BsWnD/mvp/LcJ/CJ7+/yEsC00jeI/ymfG+APy+
e+9voIU6KhEreAfJ8g==
-----END CERTIFICATE-----
```

Set the new credential in kubeconfig for administrators to use

- `kubectl config set-credentials — client-key=jon.key — client-certificate=jon.crt — embed-certs`
- `kubectl config view`

[Open in app](#)

```

apiVersion: v1
clusters:
- cluster:
  certificate-authority-data: DATA+OMITTED
  server: https://192.168.1.4:6443
  name: kubernetes
contexts:
- context:
  cluster: kubernetes
  user: kubernetes-admin
  name: kubernetes-admin@kubernetes
current-context: kubernetes-admin@kubernetes
kind: Config
preferences: {}
users:
- name: jon
  user:
    client-certificate-data: REDACTED
    client-key-data: REDACTED
- name: kubernetes-admin
  user:
    client-certificate-data: REDACTED
    client-key-data: REDACTED

```

Set new context in the cluster

- `kubectcl config set-context jon — user=jon — cluster=kubernetes`

Use the context

- `kubectcl config use-context jon`

```

jonw@CKS-Master:~$ kubectcl config set-context jon --user=jon --cluster=kubernetes
Context "jon" created.
jonw@CKS-Master:~$ kubectcl config use-context jon
Switched to context "jon".

```

Depending on what permissions have been given to user “jon”, the machine could perform different actions on the credential provided.

Service Account in Pods

The topic is to ensure administrators are not giving service accounts within Pods to have permissions besides required. For demonstration, we would be creating a service account

- `kubectcl create sa podsa`

and we would see a secret is also auto generated associated with the service account

[Open in app](#)

```
jonw@CKS-Master:~$ kubectl create sa podsa
serviceaccount/podsa created
jonw@CKS-Master:~$ kubectl get secrets
```

NAME	TYPE	DATA	AGE
default-token-b5fkr	kubernetes.io/service-account-token	3	21h
podsa-token-q6xbf	kubernetes.io/service-account-token	3	8s
secure-ingress	kubernetes.io/tls	2	3h36m

Then, create a Pod that uses that service account and allow the Pod to use its auto-generated service account token (If not explicitly denied, the default is always allow.)

```
jonw@CKS-Master:~$ cat pod-usesa.yaml
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    run: usesa
  name: usesa
spec:
  serviceAccountName: podsa
  automountServiceAccountToken: true
  containers:
  - image: nginx
    name: usesa
    resources: {}
  dnsPolicy: ClusterFirst
  restartPolicy: Always
status: {}
```

When we execute inside a shell of the Pod, we could see the auto-generated service account token.

- `kubectl exec usesa -it — bash`
- `mount | grep sec`
- `cd /run/secrets/kubernetes.io/serviceaccount`
- `cat token`

```
root@usesa:/# mount | grep sec
tmpfs on /run/secrets/kubernetes.io/serviceaccount type tmpfs (ro,relatime)
root@usesa:/# cd /run/secrets/kubernetes.io/serviceaccount
root@usesa:/run/secrets/kubernetes.io/serviceaccount# ls
ca.crt namespace token
root@usesa:/run/secrets/kubernetes.io/serviceaccount# cat token
```

[Open in app](#)

```
lcnZpY2VhY2NvdW50OmRLZmF1bHQ6cG9kc2EifQ.cq90jLhDi3AFqJGCMcitCiVVYrto6UkyCs7LHXi2chxkJ5eIsJ08SldZ4YqdtU
liKgDpAl9B6-eJgJfLWPoWDSG7dHgC7NwL6G6aLq-15hG4hQ40sDjr5EKK1qLOR5UfkIrpGtnb2YveLfcq6E0FwgV3dTs7kQfpo_yL
2qfvUvA7wZWduOnQ2javwZu0yn0tdV2lir_pDP5osLYvqeYLBbKd8gX677l8AnJX6gITkdYRkw5e98mnlistIPkDVpZknr3mqRLsky
bq7YJdTm6OfFkMtXq827ORmtUD45L0M3HgVm-Mg-kU9ACfsNewo42yzG3lRrHg9FCKUf-_0mQsMAroot@usesa:/run/secrets/ku
```

If we recreate the Pod with `DISALLOWING` it to use the auto-generated service account token, we see no auto-generated tokens being mounted in the running container.

```
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    run: usesa
  name: usesa
spec:
  serviceAccountName: pods-a
  automountServiceAccountToken: false
  containers:
  - image: nginx
    name: usesa
    resources: {}
  dnsPolicy: ClusterFirst
  restartPolicy: Always
status: {}
```

```
jonw@CKS-Master:~$ kubectl exec usesa -it -- bash
root@usesa:/# mount | grep sec
root@usesa:/#
```

Last but not least, we could apply best practices by limit service account permissions with the right roles or cluster roles.

Kube API Server Access Management

Kube API server is considered the brain of K8s, so we would need to take extra caution what could get access to this core service. By default, kube-apiserver would allow anonymous access as we could see we get HTTP status 403 forbidden access when executing

- `curl https://localhost:6443 -k`

```
jonw@CKS-Master:~$ curl https://localhost:6443 -k
```


[Open in app](#)

```

},
"status": "Failure",
"message": "forbidden: User \"system:anonymous\" cannot get path \"/\".",
"reason": "Forbidden",
"details": {

},
"code": 403
}
jonw@CKS-Master:~$ |

```

If the request cannot even start the authentication, most likely, the console would return HTTP status 401 unauthorized according to [this section](#) of official documentation. The way to DISABLE anonymous access is by adding an additional parameter in kube-apiserver

- `sudo nano /etc/kubernetes/manifests/kube-apiserver.yaml`

```

apiVersion: v1
kind: Pod
metadata:
  annotations:
    kubeadm.kubernetes.io/kube-apiserver.advertise-address.endpoint: 192.168.1.4:6443
  creationTimestamp: null
  labels:
    component: kube-apiserver
    tier: control-plane
  name: kube-apiserver
  namespace: kube-system
spec:
  containers:
  - command:
    - kube-apiserver
    - --advertise-address=192.168.1.4
    - --allow-privileged=false
    - --anonymous-auth=false
    - --authorization-mode=Node,RBAC
    - --client-ca-file=/etc/kubernetes/pki/ca.crt
    - --enable-admission-plugins=NodeRestriction
    - --enable-bootstrap-token-auth=true
    - --etcd-cafile=/etc/kubernetes/pki/etcd/ca.crt
    - --etcd-certfile=/etc/kubernetes/pki/apiserver-etcd-client.crt

```

Once this is done, wait for kube-apiserver to restart and we could again try to access kube-apiserver with anonymous user. As expected, the console is now return HTTP status 401 as the request is not even being authenticated.

- `curl https://localhost:6443 -k`

```

jonw@CKS-Master:~$ curl https://localhost:6443 -k

```

[Open in app](#)

```
} ,
"status": "Failure",
"message": "Unauthorized",
"reason": "Unauthorized",
"code": 401
} jonw@CKS-Master:~$ |
```

Let's switch the setting back on allowing anonymous access to kube-apiserver and see how we could expose kube-apiserver service for external access.

First thing first, change service "kubernetes" to be exposed from ClusterIP to NodePort, so we could use nodes' public/private IP address and assigned port to access. Since I do not have another VM setup in the same network environment the K8s cluster is in, I would use nodes' public IP address for demonstration.

```
jonw@CKS-Master:~$ kubectl get svc
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	NodePort	10.96.0.1	<none>	443:30840/TCP	20h
service1	ClusterIP	10.106.255.37	<none>	80/TCP	3h33m
service2	ClusterIP	10.111.145.107	<none>	80/TCP	3h32m

Do a simple curl test to see whether we are getting HTTP status 403

- `curl https://52.137.121.234:30840 -k`

```
jonw@CKS-Master:~$ curl https://52.137.121.234:30840 -k
{
  "kind": "Status",
  "apiVersion": "v1",
  "metadata": {
  },
  "status": "Failure",
  "message": "forbidden: User \"system:anonymous\" cannot get path \"/\"",
  "reason": "Forbidden",
  "details": {
  },
  "code": 403
}
```

Now, we ensure we are using kube-apiserver-recognized FQDN or IP address

Open in app



```

jonw@CKS-Master:~$ openssl x509 -in /etc/kubernetes/pki/apiserver.crt -text
Certificate:
  Data:
    Version: 3 (0x2)
    Serial Number: 5088853778678635393 (0x509d31bc3a3f1d81)
    Signature Algorithm: sha256WithRSAEncryption
    Issuer: CN = kubernetes
    Validity
      Not Before: Feb  2 23:43:28 2021 GMT
      Not After : Feb  2 23:43:28 2022 GMT
    Subject: CN = kube-apiserver
    Subject Public Key Info:
      Public Key Algorithm: rsaEncryption
      RSA Public-Key: (2048 bit)
      Modulus:
        08:bf:95:9e:bd:5d:b3:48:07:56:93:87:f9:4c:ed:
        5a:78:c2:9a:3b:59:23:29:e8:d6:be:61:da:8e:3b:
        b6:52:6d:05:df:80:3a:85:7d:01:41:48:04:b7:6b:
        54:07:dd:75:b4:ca:1f:bc:af:49:36:07:74:8e:57:
        48:cb:28:3c:a6:cc:94:14:f5:26:80:5e:97:0b:9e:
        54:1c:b8:75:f6:23:27:2f:90:3e:9e:c4:d5:01:a4:
        7f:c1:95:e1:d5:3d:e5:11:70:4d:b2:45:2d:e7:c8:
        c2:64:cf:8b:ab:56:bc:00:a6:86:be:6b:cc:68:c6:
        57:b1:59:00:03:f3:72:fe:6e:f5:35:ec:aa:7a:0b:
        24:cb:ce:9d:e3:1b:63:9c:02:6d:cf:1e:f5:66:b6:
        5a:de:e5:43:96:ca:72:c3:f3:f8:b0:24:c7:8a:23:
        b1:02:44:24:3e:da:12:cd:2e:9a:1d:83:bf:ed:24:
        84:01:d5:e7:04:c7:0c:cb:48:7d:e5:80:59:7a:49:
        99:94:84:36:61:01:79:73:fa:24:47:1f:16:46:f2:
        c8:0a:2b:47:b8:4d:1e:25:c2:63:23:c0:f6:70:13:
        5c:e8:4b:2a:b7:09:02:32:e8:de:0f:2b:39:c3:be:
        71:ef:a1:d5:0f:af:8f:3f:07:00:7e:5e:e7:bd:0b:
        ec:25
      Exponent: 65537 (0x10001)
    X509v3 extensions:
      X509v3 Key Usage: critical
        Digital Signature, Key Encipherment
      X509v3 Extended Key Usage:
        TLS Web Server Authentication
      X509v3 Authority Key Identifier:
        keyid:1E:E9:ED:46:AD:11:F4:91:C8:24:9C:03:C6:34:7E:5A:9B:04:50:D6
      X509v3 Subject Alternative Name:
        DNS:kubernetes-master, DNS:kubernetes, DNS:kubernetes.default, DNS:kubernetes.default.svc, DNS:kubernetes.default.svc.cl
        uster.local, IP Address:10.96.0.1, IP Address:192.168.1.4
  
```

Edit hosts to let any of the FQDN we are seeing in the image above to be resolved in nodes' public IP address required (nodes' public/private IP address).

- `sudo nano /etc/hosts`

Look up the kubernetes and see what it resolves into

- `nslookup kubernetes`

```

jonw@CKS-Master:~$ cat /etc/hosts
127.0.0.1 localhost
52.183.126.105 kubernetes
# The following lines are desirable for IPv6 capable hosts
::1 ip6-localhost ip6-loopback
fe00::0 ip6-localnet
ff00::0 ip6-mcastprefix
ff02::1 ip6-allnodes
ff02::2 ip6-allrouters
ff02::3 ip6-allhosts
jonw@CKS-Master:~$ nslookup kubernetes
Server:      127.0.0.53
  
```




Open in app

```
Name:      Kubernetes
Address: 52.183.126.105
```

View the raw data of kubectl config content.

- *kubectrl config view* — *raw > conf*

```

john@CKS-Master:~$ kubectl config view --raw
apiVersion: v1
clusters:
- cluster:
    certificate-authority-data: LS0tLS1CRUdJTI8DRVJUSUZJQ0FURSB0tLS0tCk1JSUM1ekNDQWMrZ0F3SUJBZ0lCQURBTKJna3Foa2lHOXc
wQkFRc0ZBREFFWTVJnd0VRWURWUVERXdwcmRXSmmwKY20lBGRHVnpNQjRYRFRJeE1ESXhOREF3TVRVVeE1Wb1hEVE14TURJeE1QXQdMVFV4TVZvd0ZURV
RnQkVHQFTVRQpBeE1LYTNWavPYSnVaWfJ5y3pDQ0FTSXdeUUVlKS29aSwH2Y05BUUVCQLFBRGdnRvBBRENDQVFvQ2dnRUJBSjducClFkdStDSVJ1cjdyT
FQwa19jSFhmQXZXRdVal0ZXdVpyZmh4a0hTdU01NTiXVksZl2FGe1LNRGVadeJVCkREdXMKT3RVRjBxMXNTT1pjR2V0Zk5uTGvYU3UzZlRlc29JUTJz
SKVhWGZgUmpnK3pxdXhZc2pqTjFiQjZvMlUvNVRTQpp5SR3F5a1N0Q210cXhoT1RPVVBCTEZR2hHskJzWHRHT0t5TndmTUVK0UMvRy9PQW9HaHpVYmY
vNVLfVksHYmtzCmRSu2M1WUhmD093L2tXm1ZaWGXuXNJRk1DZnZ2VjBmcGtCwNVfTk1TRzhrYVNIU2hVUNNR25sUFJodUlKOHUKemS5ZEVENovJ
hxamZorWxXd1B1M0JIWwdNc1SVfD0rLBUT0cxVG5j23pRTTFC0EtqbTb6dlNBL0FSZHBibwpoYmmlakxJVTdpWwFGL1JGcC9jQ0F3RUFBU5DTUvBd
0RnWURWjLQ0ZVFtL0ZBUURBZ0trTUEAR0E5VWRFd0VCc193UuZnQ1lCQWY4d0hRWURWujBpQkJZRUZJbk5TYUFvN1R1QzI2iWjYzemSVN1hDSVBvMdmD
QTBHQ1NXR1NJYjMkRFFQkN3VUFBNELCQVFBY2lTr3hLcUFhYm1Gd09SU3lUdElhejk5bHNmZ0lUMXBwZVFML2R2cUdva1lnZmdaTwo5WUTkKc0yYUL
aUW5zUUDJWwG0YURMY3B4akFxbEtZwnkxamtrBstHTlhoYtdyOU53VDFZVCs5RXVNC3px5nJkCnLZK3RMEvhpVG1EYl05qZIrQ3p2N0oyNU9ycUN5dz
htL3pRSG1JY1REMF0ZHGm1L13ZEULI4VXNlVTJENW0KMKZqewXQRUXdHIZZTCvZWNNWwNkM0NacWEZMDVHYTC4YUppnnNECUREaGt0R0VKM0xue
HB5MlXqSNWGDND0Uwp4b1BqL21FTJZueTFxVUVGSwtWm4wKkLSZFV3WEZHblc1RFLBREpNVdNdp2Zhe1pvWnVlUj1BZ3A1a2I4ME1ScnLLaytacWkz
YXNpWlQdTN0gZewpaUjEwTmRQ93VhNRf1ow9ZBq0tLS0tLVORCDBRVJUSUZJQ0FURSB0tLS0tCg==
    server: https://192.168.1.4:6443
  name: kubernetes
contexts:
- context:
    cluster: kubernetes
    user: kubernetes-admin
    name: kubernetes-admin@kubernetes
current-context: kubernetes-admin@kubernetes
kind: Config
preferences: {}
users:
- name: kubernetes-admin
  user:
    certificate-authority-data: LS0tLS1CRUdJTI8DRVJUSUZJQ0FURSB0tLS0tCk1JSURFekNDQWZlZ0F3SUJBZ0lJRkdnaFhIVkhkTjh3RFFZSK
tvWkLodmNOQVFFTEJRQXdGVEVUTUJFR0ExVUUKQXhNS2EzVmlaWEp1WlshbSGN6QWVGdZB5TVRBeU1UUXdNREUxTVRGYUz3MhLNAkF5TVRRd01ERTFNv
E5hTURREAgCekFWQmd0VkJB1REBbk41YzNSBgJCCHRZWE4wWlHkE15a3dGd1ELVFRREV4QnJkV0psY20lBGRHVnpMV0ZrCnJXBHVN5LCSWpBTk3n
a3Foa2lHOXc0ekFRRUZBUURBZ0R0VE4QU1JSUJ0Z0t0QVFFQWw0N1LVWGUuQ0EtH0VHb0kKcXQwC043cTlAZmJ3NjZ5NFRHZEVRWnJrNjhkT0JjcFhleVp
EcEowZGg5a0Nnc1RjvVWMrZUM3azkyRzRGt0tLMApOTWhC12Z0G1GWxdKbG9hTG50NHNpbWl0TGRYQ29EYWMZdnZHM1ZlTLJRM03XamdBUg51UnFJWU
FhVDJxMDFLC1E3emthRz1ic1dne1BtH0k4TstCa3E1aFhLEURjd2FJR056bHM0ZWZ5dDjQ2p0R21YN2t5N1pIMGk3S1taY08KQ3BxdK9rMTHwbW1hN
TR3YjZtaG5wSmlwTE5BYTdLNGxtZvYz1B31Rz1MSEXSQnLRsk43WEs1TXZFSjvYQZFTcgpyVlRvZUZZn15QmhyNXd2T1BwN1VscXN3SDV0bzLDWVZq
nT3ZMkNz0etWRLg4tzhvSE9TUjF0bCsYU5V6M31vCms0UmRd0LEQVFBQm8wZ3dsakFPQnd0VkhROEJBZjhfQkFNQ0JhQXdDf1LEVLiWbEJBd3dDZ1L
JS3dZQkJRVUgHQXZjd0h3WURWUjBqQkKjnd0ZvQVVPYzFKb0NqdE9NR1Z0bnJmT2RUDGNJZytsV0F3RFFZSktvWkLodmNOQVFFTApcUUFEEZ2dFQkFCcz
L2dGZSY5S9SbFhZMTFsVE0Yzk9pNEVZQ0pNcnM3aUdIMk1J5TThJVkhRc0drNWJlQmVvT1LHCm1BWw04Q0t0YTZlUWY3aCtnWUY1M3p1M3phUGZLD0x8c
mFISF5aU1Z3AHNiVlR3dWxdZnA5bG0vUjJ4QW54MmUKVVR5wi9a51ZhTmdoQZq5cXZCaxZuN1Z3U1H3bVvQTTd3VHFvBfDZNEk4dnVxVEUwVWtDL2FC
N0uSgamlTaDM3cQpQ3VdWY4YU1saGdJbXB5SRT1IUXDp2FMhR5WELCVUFDZ2XWk9MCFG2d041VmpWkRrNWMEHYVfISlhaNnBLcULcMlpQh3XamVjT3Z
kd1lvZzd0V3JtQ0FRnmUrY1JGU1ZvelZLdjRnckluLzZVWVkh0Kd5UVNuzWi90Z1NCVDdPL2UKd3pZTLArbjN0K1RvR2paandwOVBSZLDQWjIrRFFiND
0KLS0tLS1FTkQ0QV5SVELGSUNBVEU1S0tLQ0=
  client-key-data: LS0tLS1CRUdJTI8SOU0EUFJJVkJFURSB1RvktLS0tLQpNSU1Fb3dJQkFBS0NB0UUVBbDQ3WVYan5DS0dTRudvSXFOHMHNON3
E5WmZidzY2eTRURZFa1pyaz14ZE9CY3BYCmV5KRW5jBkaDlRq2dyVGVNVUytLQzdrOTJHNEZPS0SwTk1oQnMWWThTR1L3SnxvYUxuaDRzaW1INEkKw
ENvRGEKYzN2dkczVkt0ULeZqlQdZ0FQazVScULZQWUMnEmWVRN3prYUcyNXXNZ3pQbTNJOE0RQmtXNWHYS31EY3dhYwPhtKZsczRLYzZIN0LDak5H
bvG3a3k3WkgwaTdKK1pJ70NwCZPazE4cG1tYTU0d2YbWhucELpcE0XQW5Z3TRrCnFPNXIvcHVH0UxITFJCeVFvTjdySvVndkVKNVhDMVnycLZub2V
GcZ25eUJocjV3dk9KcQdvbfHfZd0g1Tm85Q1kKvMoxNzMyQ3NXS01GWDhPOG5IT1NSMhR5KsytVTm8zci9rNFJja3dJREFRQJJBb0LCQh0r0WlXanJpZz
ZLwZ5ZaPpWjKfXNFCYj3mV1EvWY4U1saGdJbXB5SRT1IUXDp2FMhR5WELCVUFDZ2XWk9MCFG2d041VmpWkRrNWMEHYVfISlhaNnBLcULcMlpQh3XamVjT3Z
UpWalaowazBFYKxocC9NRDh6d3cx0VNGTX0e2ZkE2UXFXdUS2d3dZefJSR1ZrL2FTQ3kkkWE1WZGQyN2UWQnpPbWfM5VZXWg02ZVc0WUNTc2x1UHXn0Gpx
ewdyUURYZHo4d59BcFAWm2tsdkVqaTN3amtaNw0RDNuK1LaRndDYU00Smo5QWsyQWm0GdFZTJv5whUHLYSWyaFABET19PbLRnBERCMkVZUkRSZWo

```

Save the existing raw CA certificate, client certificate and key information into file “conf” and modify the server information to use nodes’ public/private IP address or in this case FQDN if you have configured in hosts record.

```
apiVersion: v1
clusters:
```

[Open in app](#)

```
- context:
  cluster: kubernetes
  user: kubernetes-admin
  name: kubernetes-admin@kubernetes
current-context: kubernetes-admin@kubernetes
kind: Config
preferences: {}
users:
- name: kubernetes-admin
  user:
    client-certificate-data: LS0tLS1CRUdJTiB0RVJUSUZJQ0FURSB0tLS0tCk1JSURFekNDQWZ1Z0F3SUJBZ0lJRkdnaFhIVkhkTjh3RFFZS$
    client-key-data: LS0tLS1CRUdJTiB0RVJUSUZJQ0FURSB0tLS0tCk1JSURFekNDQWZ1Z0F3SUJBZ0lJRkdnaFhIVkhkTjh3RFFZS$
```

Try contacting kube-apiserver, try to get namespaces in this case, with the modified conf file.

- `kubect — kubeconfig conf get ns`

```
jonw@CKS-Master:~$ kubectl --kubeconfig conf get ns
NAME                STATUS    AGE
default             Active    20h
ingress-nginx       Active    3h40m
kube-node-lease     Active    20h
kube-public         Active    20h
kube-system         Active    20h
kubernetes-dashboard Active    18h
```

Node Restriction

Check on K8s master nodes and see whether NodeRestriction admission plug-in is enabled already

- `sudo cat /etc/kubernetes/manifests/kube-apiserver.yaml`

```
apiVersion: v1
kind: Pod
metadata:
  annotations:
    kubeadm.kubernetes.io/kube-apiserver.advertise-address.endpoint: 192.168.1.4:6443
  creationTimestamp: null
  labels:
    component: kube-apiserver
    tier: control-plane
  name: kube-apiserver
  namespace: kube-system
spec:
  containers:
  - command:
    - kube-apiserver
    - --advertise-address=192.168.1.4
    - --allow-privileged=false
    - --anonymous-auth=true
    - --authorization-mode=Node,RBAC
```


[Open in app](#)

```
- --etcd-certfile=/etc/kubernetes/pki/apiserver-etcd-client.crt
- --etcd-keyfile=/etc/kubernetes/pki/apiserver-etcd-client.key
```

Head over to K8s worker nodes and check whether using worker nodes' kubelet context could modify master nodes labels. First, we would need to associate worker nodes' kubelet context as default context for executing K8s CLI

- `sudo su`
- `export KUBECONFIG=/etc/kubernetes/kubelet.conf`

and test with whatever command works or not

- `kubectl get ns`

As expected, worker nodes' kubelet does not have permission to get namespaces.

```
root@CKS-Worker:/home/jonw# export KUBECONFIG=/etc/kubernetes/kubelet.conf
root@CKS-Worker:/home/jonw# kubectl get ns
Error from server (Forbidden): namespaces is forbidden: User "system:node:cks-worker" cannot list resource "namespaces" in API group "" at the cluster scope
```

With worker nodes' kubelet context, it is not authorized to label master nodes but it is able to label worker nodes (itself or other nodes).

- `kubectl label node cks-master cks/test=yes`
- `kubectl label node cks-worker cks/test=yes`

```
root@CKS-Worker:/home/jonw# kubectl label node cks-master cks/test=yes
Error from server (Forbidden): nodes "cks-master" is forbidden: node "cks-worker" is not allowed to modify node "cks-master"
root@CKS-Worker:/home/jonw# kubectl label node cks-worker cks/test=yes
node/cks-worker labeled
```

Upgrade Kubernetes Clusters with kubectl

One of the most common tasks every IT administrator would need to do is to update or upgrade the running machines. For K8s administrators, K8s version would also need to be maintain in supported scope and the following would show how to do it.

Master Nodes

Get all the apt updates

[Open in app](#)

Make sure no more pods being scheduled to master nodes.

- `kubectcl cordon cks-master`

```
jonw@CKS-Master:~/container$ kubectcl cordon cks-master
node/cks-master cordoned
jonw@CKS-Master:~/container$ kubectcl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
cks-master	Ready, SchedulingDisabled	control-plane, master	5d	v1.20.2
cks-worker	Ready	<none>	5d	v1.20.2

Drain all pods, deployment from master nodes.

- `kubectcl drain cks-master — ignore-daemonsets`

Check kubeadm version

- `kubeadm version`

```
jonw@CKS-Master:~/container$ kubeadm version
kubeadm version: &version.Info{Major:"1", Minor:"20", GitVersion:"v1.20.2",
```

Get the upgrade plan

- `kubeadm upgrade plan`

```
jonw@CKS-Master:~/container$ sudo kubeadm upgrade plan
[sudo] password for jonw:
[upgrade/config] Making sure the configuration is correct:
[upgrade/config] Reading configuration from the cluster...
[upgrade/config] FYI: You can look at this config file with 'kubectcl -n kube-system get cm kubeadm-config -o yaml'
[preflight] Running pre-flight checks.
[upgrade] Running cluster health checks
[upgrade] Fetching available versions to upgrade to
[upgrade/versions] Cluster version: v1.20.2
[upgrade/versions] kubeadm version: v1.20.2
[upgrade/versions] Latest stable version: v1.20.2
[upgrade/versions] Latest stable version: v1.20.2
[upgrade/versions] Latest version in the v1.20 series: v1.20.2
[upgrade/versions] Latest version in the v1.20 series: v1.20.2
```

Apply the upgrade plan shown before

- `kubeadm upgrade apply <K8s version>`

Check kubectcl and kubelet version

[Open in app](#)

```
jonw@CKS-Master:~/container$ kubectl version
Client Version: version.Info{Major:"1", Minor:"20", GitVersion:"v1.20.2",
Server Version: version.Info{Major:"1", Minor:"20", GitVersion:"v1.20.2",
```

- `kubectl get nodes -o yaml | grep kubelet`

```
jonw@CKS-Master:~/container$ kubectl get nodes -o yaml | grep kubelet
      f:kubeletEndpoint:
      f:kubeletVersion: {}
    manager: kubelet
    message: kubelet has sufficient memory available
    message: kubelet has no disk pressure
    message: kubelet has sufficient PID available
    message: kubelet is posting ready status. AppArmor enabled
    kubeletEndpoint:
    kubeletVersion: v1.20.2
      f:kubeletEndpoint:
      f:kubeletVersion: {}
    manager: kubelet
    message: kubelet has sufficient memory available
    message: kubelet has no disk pressure
    message: kubelet has sufficient PID available
    message: kubelet is posting ready status. AppArmor enabled
    kubeletEndpoint:
    kubeletVersion: v1.20.2
```

Install all core components to the required version

- `apt-get install kubeadm=<K8s version> kubelet=<K8s version> kubectl=<K8s version>`

Make master nodes available for pod scheduling once again.

- `kubectl uncordon cks-master`

Worker Nodes

Worker nodes upgrade is slightly different, but mostly the same concept.

Get all the apt updates

- `sudo apt-get update`

Make sure no more pods being scheduled to master nodes.

[Open in app](#)

Drain all pods, deployment from master nodes.

- *kubectldrain cks-worker — ignore-daemonsets*

Check kubeadm version

- *kubeadm version*

Install the required kubeadm version

- *apt-get install kubeadm = <required version>*

Upgrade worker nodes with kubeadm

- *kubeadm upgrade node*

Install all core components to the required version

- *apt-get install kubelet = <K8s version> kubectld = <K8s version>*

Check kubelet and kubectld version to ensure they are running in the required version.

- *kubelet — version*
- *kubectld — version*

Make worker nodes available for pod scheduling once again.

- *kubectlduncordon cks-worker*

For more details on how to use kubeadm to upgrade master nodes and worker nodes, please check [this site](#).

[Open in app](#)



[About](#) [Help](#) [Legal](#)

Get the Medium app

