

[Open in app](#)**Jonathan**[Follow](#)

29 Followers

[About](#)

Certified Kubernetes Security Specialist (CKS) Preparation Part 5 — Microservice Vulnerabilities



Jonathan Mar 3 · 7 min read

If you have not yet checked the previous parts of this series, please go ahead and check [Part1](#), [Part2](#), [Part3](#) and [Part4](#).

In this article, I would focus on the preparation around **microservice vulnerabilities** in CKS certification exam.

Manage Secrets

For knowing how Kubernetes store secrets and how Pods use secrets, we first create 2 secrets (secret1 and secret2).

- `kubectl create secret generic secret1 —from-literal=username=jonw`
- `kubectl create secret generic secret2 —from-literal=password=12345678`
- `kubectl get secrets`

```
jonw@CKS-Master:~$ kubectl get secrets
NAME                                TYPE                                DATA  AGE
default-token-xd7dp                kubernetes.io/service-account-token 3      46h
podsa-token-bmkk9                  kubernetes.io/service-account-token 3      46h
secret1                             Opaque                               1      19h
secret2                             Opaque                               1      19h
```

[Open in app](#)

- `kubectrl run mountsecrets — image=nginx -o yaml — dry-run=client > pod-mountsecrets.yaml`
- `nano pod-mountsecrets.yaml`
- `kubectrl create -f pod-mountsecrets.yaml`

```
jonw@CKS-Master:~$ cat pod-mountsecrets.yaml
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    run: mountsecrets
  name: mountsecrets
spec:
  containers:
  - image: nginx
    name: mountsecrets
    resources: {}
    env:
      - name: secret2
        valueFrom:
          secretKeyRef:
            name: secret2
            key: pass
    volumeMounts:
      - name: secret1
        mountPath: "/etc/secret1"
        readOnly: true
  volumes:
  - name: secret1
    secret:
      secretName: secret1
  dnsPolicy: ClusterFirst
  restartPolicy: Always
status: {}
```

After the pod gets created, we could execute into a shell inside of the Pod and try to read the secrets.

- `kubectrl exec mountsecrets -it — bash`
- `cat /etc/secret1/user`



```
root@mountsecrets:/# cat /etc/secret1/user
adminroot@mountsecrets:/#
root@mountsecrets:/# env | grep secret2
secret2=12345678
```

The permission each Pod leveraging is the default service account created in the namespace. In this example, it would be service account “default” in the namespace “default”. Each service account would automatically generate a token, the name would be similar to “<service account name>-token-<random-generated-string>”.

- `kubectl get sa`
- `kubectl get secrets`

```

jonw@CKS-Master:~$ kubectl get sa
NAME          SECRETS  AGE
default       1        4d20h
podsa         1        4d20h
jonw@CKS-Master:~$ kubectl get secrets
NAME                                TYPE                                DATA  AGE
default-token-xd7dp                 kubernetes.io/service-account-token 3      4d20h
podsa-token-bmkk9                   kubernetes.io/service-account-token 3      4d20h
secret1                             Opaque                               1      3d17h
secret2                             Opaque                               1      3d17h
secret3                             Opaque                               1      3d16h

```

If executing into one of the shell within Pods(containers), the token would be shown under file mount.

- `mount | grep serviceaccount`
- `cat /run/secrets/kubernetes.io/serviceaccount/token`

[illegible]

If any endpoint needs to be access with this token, this could be easily done with



Then, we execute ETCD CLI with the queried certificates and keys above to get the secrets are showing in plain format.

- `ETCDCTL_API=3 etcdctl --cacert="/etc/kubernetes/pki/etcd/ca.crt" --cert="/etc/kubernetes/pki/apiserver-etcd-client.crt" --key="/etc/kubernetes/pki/apiserver-etcd-client.key" get /registry/secrets/<namespace>/<secret name>`

[illegible]

By now, we understand it would be “safer” if we could have ETCD encrypted even at rest and the way to do that is by enabling it in kube-apiserver configuration file. First, we create an ETCD encryption configuration file somewhere under /etc/kubernetes. In this case, we would create a file named “ec.yaml” in a folder named “etcd” under that path. For the secret section, please use base 64 encoded form of whatever string provided. Please check [here](#) for more ETCD encryption information.

- `cd /etc/kubernetes`
- `mkdir etcd`
- `nano etcd/ec.yaml`

```
jonw@CKS-Master:~$ cat /etc/kubernetes/etcd/ec.yaml
apiVersion: apiserver.config.k8s.io/v1
kind: EncryptionConfiguration
resources:
  - resources:
    - secrets
    providers:
    - aescbc:
        keys:
        - name: key1
          secret: cGEzc3dvcnRwYXNzd29vZA==
```

[Open in app](#)

After creating the ETCD configuration file, head over to configure kube-apiserver file for using the newly created configuration file. Also, remember to add a hostVolume and volumeMount in the bottom part of the configuration file.

- `nano /etc/kubernetes/manifests/kube-apiserver.yaml`

```
apiVersion: v1
kind: Pod
metadata:
  annotations:
    kubeadm.kubernetes.io/kube-apiserver.advertise-address.endpoint: 192.168.1.4:6443
  creationTimestamp: null
  labels:
    component: kube-apiserver
    tier: control-plane
  name: kube-apiserver
  namespace: kube-system
spec:
  containers:
  - command:
    - kube-apiserver
    - --encryption-provider-config=/etc/kubernetes/etcd/ec.yaml
    - --advertise-address=192.168.1.4
    - --allow-privileged=true
    - --authorization-mode=Node,RBAC
    - --client-ca-file=/etc/kubernetes/pki/ca.crt
    - --enable-admission-plugins=NodeRestriction
    - --enable-bootstrap-token-auth=true
    - --etcd-cafile=/etc/kubernetes/pki/etcd/ca.crt
    - --etcd-certfile=/etc/kubernetes/pki/apiserver-etcd-client.crt
    - --etcd-keyfile=/etc/kubernetes/pki/apiserver-etcd-client.key
    - --etcd-servers=https://127.0.0.1:2379
```

```
volumeMounts:
  - mountPath: /etc/ssl/certs
    name: ca-certs
    readOnly: true
  - mountPath: /etc/ca-certificates
    name: etc-ca-certificates
    readOnly: true
  - mountPath: /etc/kubernetes/pki
    name: k8s-certs
    readOnly: true
  - mountPath: /usr/local/share/ca-certificates
    name: usr-local-share-ca-certificates
    readOnly: true
  - mountPath: /usr/share/ca-certificates
    name: usr-share-ca-certificates
    readOnly: true
  - mountPath: /etc/kubernetes/etcd
    name: k8s-etcd
    readOnly: true
```

```
- hostPath:
  path: /etc/ssl/certs
  type: DirectoryOrCreate
  name: ca-certs
- hostPath:
  path: /etc/ca-certificates
  type: DirectoryOrCreate
  name: etc-ca-certificates
- hostPath:
  path: /etc/kubernetes/pki
  type: DirectoryOrCreate
  name: k8s-certs
- hostPath:
  path: /usr/local/share/ca-certificates
  type: DirectoryOrCreate
  name: usr-local-share-ca-certificates
- hostPath:
  path: /usr/share/ca-certificates
  type: DirectoryOrCreate
  name: usr-share-ca-certificates
- hostPath:
  path: /etc/kubernetes/etcd
  type: DirectoryOrCreate
  name: k8s-etcd
```

Wait for the kube-apiserver to restart and test whether newly created secrets are now being encrypted even at rest. So, secret3 is created after enabling the at-rest encryption.

- `ETCDCTL_API=3 etcdctl --cacert="/etc/kubernetes/pki/etcd/ca.crt" --cert="/etc/kubernetes/pki/apiserver-etcd-client.crt" --key="/etc/kubernetes/pki/apiserver-etcd-client.key" get /registry/secrets/<namespace>/<secret name>`

```
jonn@CK8S-Master:~$ sudo ETCDCTL_API=3 etcdctl --cacert=/etc/kubernetes/pki/etcd/ca.crt --cert=/etc/kubernetes/pki/apiserver-etcd-client.crt --key=/etc/kubernetes/pki/apiserver-etcd-client.key get /registry/secrets/default/secret3/registry/secrets/default/secret3  
kBs:enc:aescbc:v1:key1::%XNNDc6+H{0xc+eD..n++6.n.DADeeDe}|F.XBWMYD.*Kf+eD++79.R.LJ.e5++Z++++?))F+++<+  
<BdVjzZigMj2+)++++++<.lnDzE8++JJQ\b+s+  
eDoHeC  
VS+++++.IeeeGGG.  
"Spae'SeeeeZeTe...Ce?eSeeee  
W++++]9ey\++  
eSdDee.,H=:DofE3uYeDe)ee5++
```

Container Runtime Interface (CRI)

[Open in app](#)

place. Before open container initiative (OCI) proposing to have CRI, the communication between containers and Kubernetes (K8s) is relying on dockershim/rkt provided and maintained by Docker. However, when containers and K8s are getting more and more sophisticated, the maintenance cost of dockershim/rkt becomes higher and higher. Therefore, having an interface that opens to the open source community and for solely dealing with container runtime becomes the answer to this challenging situation.

For CKS exam, we would need to know how to create RuntimeClass and how to create Pods using the specific RuntimeClass to communicate with K8s. For more information, please check [this site](#).

- `kubectl create -f <runtimeclass.yaml>`

```
apiVersion: node.k8s.io/v1 # RuntimeClass is defined in the node.k8s.io API group
kind: RuntimeClass
metadata:
  name: myclass # The name the RuntimeClass will be referenced by
  # RuntimeClass is a non-namespaced resource
handler: myconfiguration # The name of the corresponding CRI configuration
```

Image Credit: kubernetes.io

- `kubectl get runtimeclass`
- `kubectl create -f <pod.yaml>`

```
apiVersion: v1
kind: Pod
metadata:
  name: mypod
spec:
  runtimeClassName: myclass
  # ...
```

Image Credit: kubernetes.io

Container with Proper Permissions

[Open in app](#)

not set with proper permissions, the hosting environment are exposed to potential threat. Luckily, there are many configuration that could be implemented inside container to not running into that situation.

runAs

By adding these security contexts, whenever administrators execute into one of the Pod's terminal, the associated identity would be user:1000 and group:3000. Please check [here](#) for more information.

- `kubectl create -f <pod.yaml>`

```
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    run: runas
  name: runas
spec:
  securityContext:
    runAsUser: 1000
    runAsGroup: 3000
  containers:
  - command:
    - sh
    - -c
    - sleep 1d
    image: busybox
    name: runas
    resources: {}
  dnsPolicy: ClusterFirst
  restartPolicy: Always
  status: {}
```

- `kubectl exec <pod name> -it -- sh`
- `id`

```
jonw@CKS-Master:~$ kubectl exec runas -it -- sh
/ $ id
uid=1000 gid=3000
```

runAsNonRoot

[Open in app](#)


it would need root access to run with the image. Please check [here](#) for more information.

- `kubect create -f <pod.yaml>`

```
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    run: runasnonroot
  name: runasnonroot
spec:
  containers:
  - command:
    - sh
    - c
    - sleep 1d
    image: busybox
    name: runasnonroot
    resources: {}
    securityContext:
      runAsNonRoot: true
    dnsPolicy: ClusterFirst
    restartPolicy: Always
  status: {}
```

- `kubect describe pod <pod name>`

Events:				
Type	Reason	Age	From	Message
Normal	Scheduled	7s	default-scheduler	Successfully assigned default/runasnonroot to cks-worker
Normal	Pulled	4s	kubelet	Successfully pulled image "busybox" in 1.311995533s
Normal	SandboxChanged	4s	kubelet	Pod sandbox changed, it will be killed and re-created.
Normal	Pulling	1s (x3 over 6s)	kubelet	Pulling image "busybox"
Warning	Failed	1s (x2 over 4s)	kubelet	Error: container has runAsNonRoot and image will run as root
Normal	Pulled	1s	kubelet	Successfully pulled image "busybox" in 1.376428132s

privileged

There is no apparent way to know whether the container is running in privileged state or not. However, this could be tested by the possibility of executing some privileged actions.

- `kubect create -f <pod.yaml>`

```
apiVersion: v1
```

[Open in app](#)

```

labels:
  run: privileged
  name: privileged
spec:
  containers:
  - command:
    - sh
    - -c
    - sleep 1d
    image: busybox
    name: runas
    resources: {}
    securityContext:
      privileged: true
    dnsPolicy: ClusterFirst
    restartPolicy: Always
  status: {}

```

- `kubectl exec <pod name> -it -- sh`
- `sysctl kernel.hostname=whatever`

```

jonw@CKS-Master:~$ kubectl exec privileged -it -- sh
/ # sysctl kernel.hostname=whatever
kernel.hostname = whatever

```

allowPrivilegeEscalation

By checking `/proc/1/status` → `NoNewPrivs`, we would get the information on whether the Pod could execute privileged actions, including actions shown in the previous section “`sysctl kernel.hostname=whatever`”. If the Pod is allowed to have privilege escalation, there should be no issues.

- `kubectl create -f <pod.yaml>`

```

apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    run: allowprivilegeescalation
    name: allowprivilegeescalation
spec:
  containers:
  - command:
    - sh

```

[Open in app](#)

```
name: runas
resources: {}
securityContext:
  allowPrivilegeEscalation: true
dnsPolicy: ClusterFirst
restartPolicy: Always
status: {}
```

- `kubect exec <pod name> -it -- sh`
- `cat /proc/1/status | grep NoNewPrivs`

```
jonw@CKS-Master:~$ kubectl exec allowprivilegeescalation -it -- sh
/ # cat /proc/1/status
Name:      sleep
Umask:    0022
State:     S (sleeping)
Tgid:      1
Ngid:      0
Pid:       1
PPid:      0
TracerPid: 0
Uid:       0      0      0      0
Gid:       0      0      0      0
FDSize:    64
Groups:    10
NSTgid:    1
NSpid:     1
NSpgid:    1
NSSid:     1
VmPeak:    1300 kB
VmSize:    1300 kB
VmLck:     0 kB
VmPin:     0 kB
VmHWM:     4 kB
VmRSS:     4 kB
RssAnon:   4 kB
RssFile:   0 kB
RssShmem:  0 kB
VmData:    36 kB
VmStk:     132 kB
VmExe:     888 kB
VmLib:     0 kB
VmPTE:     28 kB
VmSwap:    0 kB
HugetlbPages: 0 kB
CoreDumping: 0
THP_enabled: 1
Threads:   1
SigQ:      1/31705
```


[Open in app](#)

```
SigIgn: 00000000000000004
SigCgt: 00000000000000000
CapInh: 00000000a80425fb
CapPrm: 00000000a80425fb
CapEff: 00000000a80425fb
CapBnd: 00000000a80425fb
CapAmb: 0000000000000000
NoNewPrivs: 0
```

PodSecurityPolicy

First, we enable kube-apiserver to apply Pod security policy within its configuration file. Then, we create one Pod security policy to determine what actions could be performed by Pod and what could not. Please check [here](#) for more information.

- `sudo nano /etc/kubernetes/manifests/kube-apiserver.yaml`

```
apiVersion: v1
kind: Pod
metadata:
  annotations:
    kubeadm.kubernetes.io/kube-apiserver.advertise-address.endpoint: 192.168.1.4:6443
  creationTimestamp: null
  labels:
    component: kube-apiserver
    tier: control-plane
  name: kube-apiserver
  namespace: kube-system
spec:
  containers:
    - command:
      - kube-apiserver
      - --encryption-provider-config=/etc/kubernetes/etcd/ec.yaml
      - --advertise-address=192.168.1.4
      - --allow-privileged=true
      - --authorization-mode=Node,RBAC
      - --client-ca-file=/etc/kubernetes/pki/ca.crt
      - --enable-admission-plugins=NodeRestriction,PodSecurityPolicy
      - --enable-bootstrap-token-auth=true
      - --etcd-cafile=/etc/kubernetes/pki/etcd/ca.crt
      - --etcd-certfile=/etc/kubernetes/pki/apiserver-etcd-client.crt
      - --etcd-keyfile=/etc/kubernetes/pki/apiserver-etcd-client.key
      - --etcd-servers=https://127.0.0.1:2379
      - --insecure-port=0
      - --kubelet-client-certificate=/etc/kubernetes/pki/apiserver-kubelet-client.crt
      - --kubelet-client-key=/etc/kubernetes/pki/apiserver-kubelet-client.key
      - --kubelet-preferred-address-types=InternalIP,ExternalIP,Hostname
```

- `kubectrl create -f psp.yaml`

```
apiVersion: policy/v1beta1
kind: PodSecurityPolicy
```

[Open in app](#)

```

spec:
  privileged: false # Don't allow privileged pods!
  allowPrivilegeEscalation: false
  # The rest fills in some required fields.
  seLinux:
    rule: RunAsAny
  supplementalGroups:
    rule: RunAsAny
  runAsUser:
    rule: RunAsAny
  fsGroup:
    rule: RunAsAny
  volumes:
    - '*'

```

When try deploying another pods allowing privilege escalation, the console would return error message because it is contradicting to Pod security policy that was applied.

- `kubectl create -f <pod.yaml>`

```

jonw@CKS-Master:~$ kubectl create -f pod-allowPrivilegeEscalation.yaml
Error from server (Forbidden): error when creating "pod-allowPrivilegeEscalation.yaml": pods "allowprivilegeescalation" is forbidden: PodSecurityPolicy: unable to admit pod: [spec.containers[0].securityContext.allowPrivilegeEscalation: Invalid value: true: Allowing privilege escalation for containers is not allowed]

```

mTLS

mTLS stands for mutual authentication, meaning client authenticates server and server does the same to client. [This Medium article](#) provides a pretty clear explanation how it works. Basically, whenever we are putting client certificate and client key in the command like “curl” or “xxxxctl”, there is a high possibility that the communication is applying mTLS. In the Medium article above, if we follow all the steps until the end, we should be having an Ingress YAML set up somewhat like below.

```

apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  annotations:
    nginx.ingress.kubernetes.io/auth-tls-verify-client: "on"
    nginx.ingress.kubernetes.io/auth-tls-secret: "default/my-certs"
  name: meow-ingress
  namespace: default
spec:
  rules:
    - host: meow.com
      http:

```

[Open in app](#)


```
servicePort: 80
path: /
tls:
- hosts:
- meow.com
secretName: my-certs
```

If we try to curl the HTTPS website without client credentials.

** The NGINX Ingress is setup to be exposed with NodePort

** The resolved public IP address could be either node's

- `curl https://meow.com:30908 --resolve meow.com:30908:52.183.91.218 -k`

```
jonw@CKS-Master:~$ curl https://meow.com:30908 --resolve meow.com:30908:52.183.91.218 -k
<html>
<head><title>403 Forbidden</title></head>
<body>
<center><h1>403 Forbidden</h1></center>
<hr><center>nginx</center>
</body>
</html>
```

If we try to curl the HTTPS website with client credentials, we would get similar results as below. We might still see the HTTP status 403 but that is just because the client is forbidden to visit the site.

- `curl https://meow.com:30908 --resolve meow.com:30908:52.183.91.218 --cert client.crt --key client.key -kv`

```
jonw@CKS-Master:~$ curl https://meow.com:30908 --resolve meow.com:30908:52.183.91.218 --cacert ca.crt --cert client.crt --key client.key -kv
* Added meow.com:30908:52.183.91.218 to DNS cache
* Rebuilt URL to: https://meow.com:30908/
* Hostname meow.com was found in DNS cache
* Trying 52.183.91.218...
* TCP_NODELAY set
* Connected to meow.com (52.183.91.218) port 30908 (#0)
* ALPN, offering h2
* ALPN, offering http/1.1
* successfully set certificate verify locations:
* CAfile: ca.crt
  CApath: /etc/ssl/certs
* TLSv1.3 (OUT), TLS handshake, Client hello (1):
* TLSv1.3 (IN), TLS handshake, Server hello (2):
* TLSv1.3 (IN), TLS Unknown, Certificate Status (22):
* TLSv1.3 (IN), TLS handshake, Unknown (8):
* TLSv1.3 (IN), TLS Unknown, Certificate Status (22):
* TLSv1.3 (IN), TLS handshake, Certificate (11):
* TLSv1.3 (IN), TLS Unknown, Certificate Status (22):
* TLSv1.3 (IN), TLS handshake, CERT verify (15):
* TLSv1.3 (IN), TLS Unknown, Certificate Status (22):
* TLSv1.3 (IN), TLS handshake, Finished (20):
* TLSv1.3 (OUT), TLS change cipher, Client hello (1):
* TLSv1.3 (OUT), TLS Unknown, Certificate Status (22):
* TLSv1.3 (OUT), TLS handshake, Finished (20):
* SSL connection using TLSv1.3 / TLS_AES_256_GCM_SHA384
* ALPN, server accepted to use h2
* Server certificate:
* subject: CN=meow.com
* start date: Feb 18 21:26:13 2021 GMT
```

Open in app



Kubernetes

Security

Preparation

Cloud Native

Microservices

About

Help

Legal

Get the Medium app

