

[Open in app](#)

Dennis Zielke

[Follow](#)

298 Followers

[About](#)

Continuous Kubernetes blue-green deployments on Azure using Nginx, AppGateway or TrafficManager — part 2



Dennis Zielke Jul 7, 2020 · 9 min read

This is part two of my series on advanced deployment practices. If you have been following [part 1](#) we finished with a working continuous deployment pipeline and some rudimentary automated rollback mechanism using helm. Unfortunately we were always replacing the existing helm release, which means that if there is something wrong with the new version, customers will always be impacted and possibly experiencing errors, while we are rolling back to the previous working version. This is obviously not ideal.

Here now enters the practice of blue/green deployments which means that instead of replacing the previous version (here we refer to this version as *blue*), we bring up the new version (here referred to as the *green* version) next to the existing version, but not expose it to the actual users right away. On the condition of having successfully validated that the green version works correctly, we will promote this version to the public version by changing the routing configuration without downtime. If something is wrong with the green version we can revert back without users every noticing interruptions.

Sound easy right? Maybe you have asked yourself the following questions:

[How to do blue/green deployments](#)

[Open in app](#)

Do you want to know how it can be done?

In this post I want to introduce you into different flavours of implementing blue/green deployments on Azure using [Nginx ingress controller](#) (any other ingress controller like [Ambassador](#), [Traefik](#) or [Kong](#) will also work), [Azure Application Gateway](#) and [Azure Traffic Manager](#) to compare their capabilities, help you make your choice to tools and provide you with samples to get you starting on adopting more advanced deployment practices in your own apps.

cloud native Prime Factor Calculator App

This is a sample application for demonstrating multi container applications on Kubernetes in Azure. [Learn more »](#)

We are going to demo green/blue deployments

Random start number

5447849

Keep posting in a loop

loop frequency 2000 ms

Request prime factors

Result:

Frontend: green-calculator-multicalculatorcanary-frontend-b9486dccb-sgh8l, green - 3.0.379

Backend: 418, green-calculator-multicalculatorcanary-backend-78c7ddc688-j9vmc, ::ffff:10.0.5.7, green - 3.0.379

Milliseconds	Value	Host	Remotelp	Version
418	[11,23,61,353]	green-calculator-multicalculatorcanary-backend-78c7ddc688-j9vmc	::ffff:10.0.5.7	green - 3.0.379
852	[67,1301,6247]	green-calculator-multicalculatorcanary-backend-78c7ddc688-j9vmc	::ffff:10.0.5.29	green - 3.0.379
1978	[b, u, g]	blue-calculator-multicalculatorcanary-backend-77fb77fd66-nnwcj		blue - 3.0.378
39	[3,139,125497]	blue-calculator-multicalculatorcanary-backend-77fb77fd66-x6cnv	::ffff:10.0.5.6	blue - 3.0.378

Users will see the blue version replaced with the green version after it has been validated — without downtime

Same as before we are aiming to achieve a couple of design considerations that we typically see in enterprise environments:

[Open in app](#)

without overhead.

2. We want to see, measure and compare telemetry and metrics for each deployment and also evaluate the blue and green deployments side by side.
3. In case of a failure we want to roll back our application automatically and have the ability to start over with the next release.

Let's get started

Again we will continue with the [phoenix](#) sample application. The easiest way to deploy everything in your own azure subscription is to open up an [azure shell](#), clone your repo there and follow these [instructions](#) to deploy at least one environment which should bring up an instance of AKS, KeyVault, Application Insights, Azure Container Registry, Application Gateway and Traffic Manager. You can check the [last post](#) on how to configure azure devops project and permissions with the newly created service principals.

The screenshot shows a GitHub commit history for the 'phoenix / .azuredavops /' branch. A single commit from 'denniszielke' is shown, committed yesterday. The commit message is partially visible as 'calculator_blue_green_build_deploy...'. Below the commit, four files are listed with their status:

File	Status
calculator_blue_green_build_deploy_appgw.yaml	added trafficmanager
calculator_blue_green_build_deploy_nginx.yaml	added trafficmanager
calculator_blue_green_build_deploy_tm.yaml	changed prefix
calculator_build_deploy.yaml	added new version

In the phoenix repository you will find an azure devops build and deployment pipeline template.

[Open in app](#)

configure the `AZURE_CONTAINER_REGISTRY_NAME` for your deployment and the `AZURE_KEYVAULT_NAME` for each environment secret store in the pipeline after you have imported it into your azure devops project.

#20200408.6 flip feature flag 1
on denniszielke.phoenix-blue-green

[Cancel](#) [:](#)

[Summary](#) [Environments](#)

Triggered by Dennis Zielke [View 2 changes](#)

Repository and version
denniszielke/phoenix
master cb0c48e

Time started and elapsed
Today at 20:44
5m 32s

Related
0 work items
1 published

Tests and coverage
[Get started](#)

[Stages](#) [Jobs](#)

Build all
1 job completed 1m 47s 1 artifact

Deploy dev1
2/3 completed 3m 31s

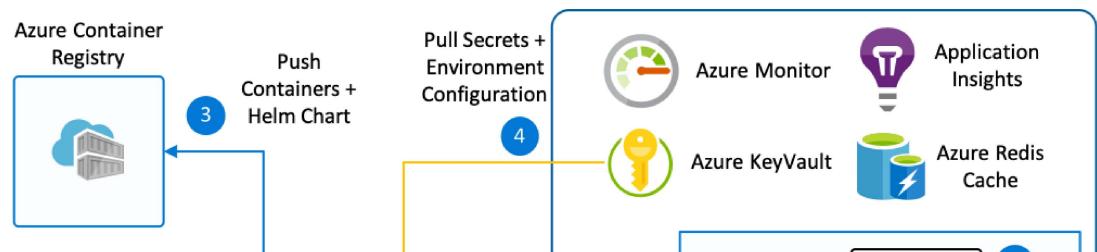
Deploy prod1
Not started

run dev1_Deploy 1m 35s
run dev1_RouteTraffic 58s
run dev1_OnSuccess 31s

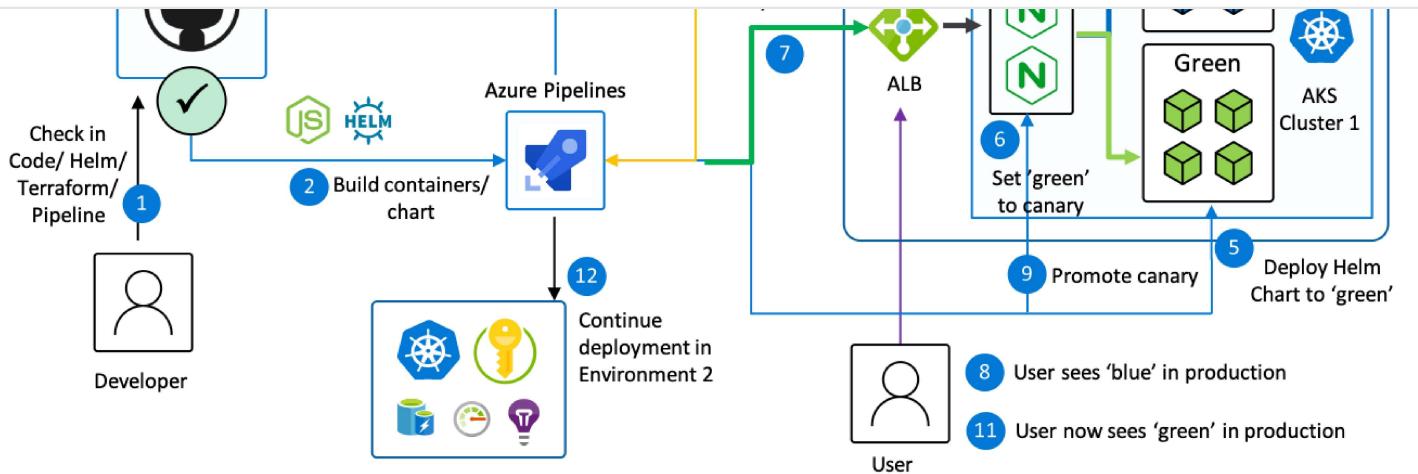
[Cancel](#)

The imported pipeline will allow you to see how the blue/green deployment progresses

We will start with the Nginx deployment, which will leverage a single azure load balancer in front of our Nginx ingress controller in the AKS cluster. The idea is to deploy our application multiple times in different namespaces leveraging the canary feature, which allows us to hook up two ingress objects to the same dns name and route them to different backend Kubernetes services depending on the existence of a custom HTTP header in the request.



[Open in app](#)



We are using the canary feature to route the traffic between the blue or the green version.

To make the automation work, we also need an extra helm variable called “canary” which will be used to define which of the two deployments is currently the canary. Before deploying a new version we will read back all deployed helm charts from the cluster, determine if there is already a canary deployment or not and perform a new deployment in either the blue or the green slot — which ever is currently not used or in a older version than the other slot.

```
scripts > deploy_multicalculator_blue_green_nginx.sh
82 if [ "$GREEN_VERSION" -gt "$BLUE_VERSION" ]; then
83     echo "Green is higher than blue - deploying blue"
84     SLOT="blue"
85 else
86     echo "Blue is higher than green - deploying green"
87     SLOT="green"
88 fi
89
90 if [ "$GREEN_VERSION" == "0" ]; then
91 if [ "$BLUE_VERSION" == "0" ]; then
92     NOCANARY="true"
93 fi
94 fi
95
96 DEPLOY_NAMESPACE=$SLOT-$KUBERNETES_NAMESPACE
97 RELEASE=$SLOT-calculator
98 kubectl create ns $DEPLOY_NAMESPACE
99
100 if [ "$NOCANARY" == "true" ]; then
101     helm upgrade $RELEASE $AZURE_CONTAINER_REGISTRY_NAME/multicalculatorcanary --namespace $DEPLOY_NAMESPACE --install --set replicaCount=4
--set image.frontendTag=$BUILD_BUILDNUMBER --set image.backendTag=$BUILD_BUILDNUMBER --set image.repository=$AZURE_CONTAINER_REGISTRY_URL
--set dependencies.useAppInsights=true --set dependencies.appInsightsSecretValue=$APPINSIGHTS_KEY --set dependencies.useAzureRedis=true
--set dependencies.redisHostValue=$REDIS_HOST --set dependencies.redisKeyValue=$REDIS_AUTH --set slot=$SLOT --set ingress.host=$INGRESS_FQDN --set ingress.class=nginx --wait --timeout 60s
102 else
103     helm upgrade $RELEASE $AZURE_CONTAINER_REGISTRY_NAME/multicalculatorcanary --namespace $DEPLOY_NAMESPACE --install --set replicaCount=2
--set image.frontendTag=$BUILD_BUILDNUMBER --set image.backendTag=$BUILD_BUILDNUMBER --set image.repository=$AZURE_CONTAINER_REGISTRY_URL
--set dependencies.useAppInsights=true --set dependencies.appInsightsSecretValue=$APPINSIGHTS_KEY --set dependencies.useAzureRedis=true
--set dependencies.redisHostValue=$REDIS_HOST --set dependencies.redisKeyValue=$REDIS_AUTH --set slot=$SLOT --set ingress.host=$INGRESS_FQDN --set ingress.class=nginx --set canary=true --set ingress.weight=0 --wait --timeout 60s
104 fi
```

The deployed helm charts in the cluster will act as source of truth — no other state store is required.


[Open in app](#)

header in the ingress object. At the same time we have defined the weight of the canary to be 0, which will ensure that no traffic will be routed there unless the header is set.

charts > multicalculatorcanary > templates > ingress.yaml > {map}

You, a few seconds ago | 2 authors (Dennis Zielke and others)

```

1 {{- if .Values.ingress.enabled -}}
2   apiVersion: networking.k8s.io/v1beta1
3   kind: Ingress
4   metadata:
5     name: {{ include "multicalculatorcanary.fullname" . }}
6     labels:
7       {{- include "multicalculatorcanary.labels" . | nindent 4 -}}
8     annotations:
9       kubernetes.io/ingress.class: {{ .Values.ingress.class }}
10    {{- if .Values.canary -}}
11      appgw.ingress.kubernetes.io/connection-draining: "true"
12      appgw.ingress.kubernetes.io/connection-draining-timeout: "60"
13      appgw.ingress.kubernetes.io/backend-path-prefix: "/"
14      nginx.ingress.kubernetes.io/canary: "true"
15      nginx.ingress.kubernetes.io/canary-by-header: "canary"
16      nginx.ingress.kubernetes.io/canary-weight: "{{ .Values.ingress.weight }}"
17    {{- with .Values.ingress.annotations -}}
18      {{- toYaml . | nindent 4 -}}
19    {{- end -}}

```

The canary variable will signal if the deployment is a canary and ensure the right ingress annotations.

As you can see below in the initial version we only have the blue version 3.0.378 deployed in the blue-calculator namespace.

NAME	NAMESPACE	REVISION	UPDATED	STATUS	CHART	APP VERSION
blue-calculator	blue-calculator	3	2020-04-08 18:41:58.31606889 +0000 UTC	deployed	multicalculatorcanary-0.1.378	3.0.378
nginx-ingress	nginx	1	2020-04-08 12:57:42.692218 +0200 CEST	deployed	nginx-ingress-1.36.0	0.30.0

Only the blue version is deployed.

Now in step 5 of our process we are deploying the green version 3.0.379 into the green-calculator namespace — but with the canary annotation in the ingress object.

NAME	NAMESPACE	REVISION	UPDATED	STATUS	CHART	APP VERSION
blue-calculator	blue-calculator	3	2020-04-08 18:41:58.31606889 +0000 UTC	deployed	multicalculatorcanary-0.1.373.0.378	3.0.379
green-calculator	green-calculator	1	2020-04-08 18:47:24.374216313 +0000 UTC	deployed	multicalculatorcanary-0.1.373.0.379	3.0.379
nginx-ingress	nginx	1	2020-04-08 12:57:42.692218 +0200 CEST	deployed	nginx-ingress-1.36.0	0.30.0

Both blue and green are deployed in the same cluster


[Open in app](#)

“canary: always” in the HTTP request we are ending up in the green version.

```
[*] ~ (⎈ dztenix-472:default) echo ""
echo "Checking production curl http://$INGRESS_FQDN/ping without canary header"
curl -s -H "canary: never" -H "Host: $INGRESS_FQDN" http://$INGRESS_FQDN/ping
echo ""
echo "Checking staging slot curl http://$INGRESS_FQDN/ping with header 'canary: always'"
curl -s -H "canary: always" -H "Host: $INGRESS_FQDN" http://$INGRESS_FQDN/ping
echo ""

Checking production curl http://51.138.7.236.xip.io/ping without canary header
{"response": "pong!", "host": "blue-calculator-multicalculatorcanary-frontend-cfc6c5b99-7xnc5", "version": "blue - 3.0.378"}
Checking staging slot curl http://51.138.7.236.xip.io/ping with header 'canary: always'
{"response": "pong!", "host": "green-calculator-multicalculatorcanary-frontend-b9486dccb-sgh81", "version": "green - 3.0.379"}
```

Compare the -H “canary:never” header value and you can see that we end up getting different versions from the /ping api

At this point we are in step 7 in our process diagram using the routeTraffic step in azure devops while using the custom headers to validate that the application works as expected. Only if that is the case we will allow the deployment to progress and end up in the on success step — if not we will execute the on failure step to clean up the canary and delete the helm deployment in the green slot.

Assuming everything worked out we will promote the canary by upgrading the helm deployment with canary=true and weight=100. This will basically override the existing ingress object for the same DNS name in the other namespace and ensure that from now on all traffic will be routed to the green deployment slot.

```
scripts > [ on_deploy_success_multicalculator_blue_green_nginx.sh
56  if [ "$CANARY_SLOT" != "none" ]; then
57    echo "Canary $CANARY_SLOT will be promoted to production"
58    DEPLOY_NAMESPACE=$CANARY_SLOT-$KUBERNETES_NAMESPACE
59    RELEASE=$CANARY_SLOT-calculator
60
61    helm upgrade $RELEASE $AZURE_CONTAINER_REGISTRY_NAME/multicalculatorcanary --namespace $DEPLOY_NAMESPACE --install --set replicaCount=2
--set image.frontendTag=$BUILD_BUILDNUMBER --set image.backendTag=$BUILD_BUILDNUMBER --set image.repository=$AZURE_CONTAINER_REGISTRY_URL
--set dependencies.useAppInsights=true --set dependencies.appInsightsSecretValue=$APPINSIGHTS_KEY --set dependencies.useAzureRedis=true
--set dependencies.redisHostValue=$REDIS_HOST --set dependencies.redisKeyValue=$REDIS_AUTH --set slot=$CANARY_SLOT --set ingress
class=nginx --set ingress.host=$INGRESS_FQDN --set canary=true --set ingress.weight=100 --wait --timeout 45s      You, a day ago * added tfi
62
63  if [ "$PRODUCTION_SLOT" != "none" ]; then
64    echo "Production $PRODUCTION_SLOT will be deleted"
65    DEPLOY_NAMESPACE=$PRODUCTION_SLOT-$KUBERNETES_NAMESPACE
66    RELEASE=$PRODUCTION_SLOT-calculator
67    helm delete $RELEASE --namespace $DEPLOY_NAMESPACE
68  fi
```

By setting canary=true and weight=100 we will override the production slot and force all traffic to the green slot.

Now we will simply delete the existing production slot and as a last step do one more upgrade of the existing green slot with canary=false to promote it to be the new production slot and ensure we can start over at the beginning when the next release with a new canary gets deployed.


[Open in app](#)

```
curl -s -H "canary: always" -H "Host: $INGRESS_FQDN" http://$INGRESS_FQDN/ping
echo ""

Checking production curl http://51.138.7.236.xip.io/ping without canary header
{"response":"pong!","host":"blue-calculator-multicalculatorcanary-frontend-cfc6c5b99-7xnc5","version":"blue - 3.0.378"}
Checking staging slot curl http://51.138.7.236.xip.io/ping with header 'canary: always'
{"response":"pong!","host":"blue-calculator-multicalculatorcanary-frontend-cfc6c5b99-gh9kf","version":"blue - 3.0.378"}
```

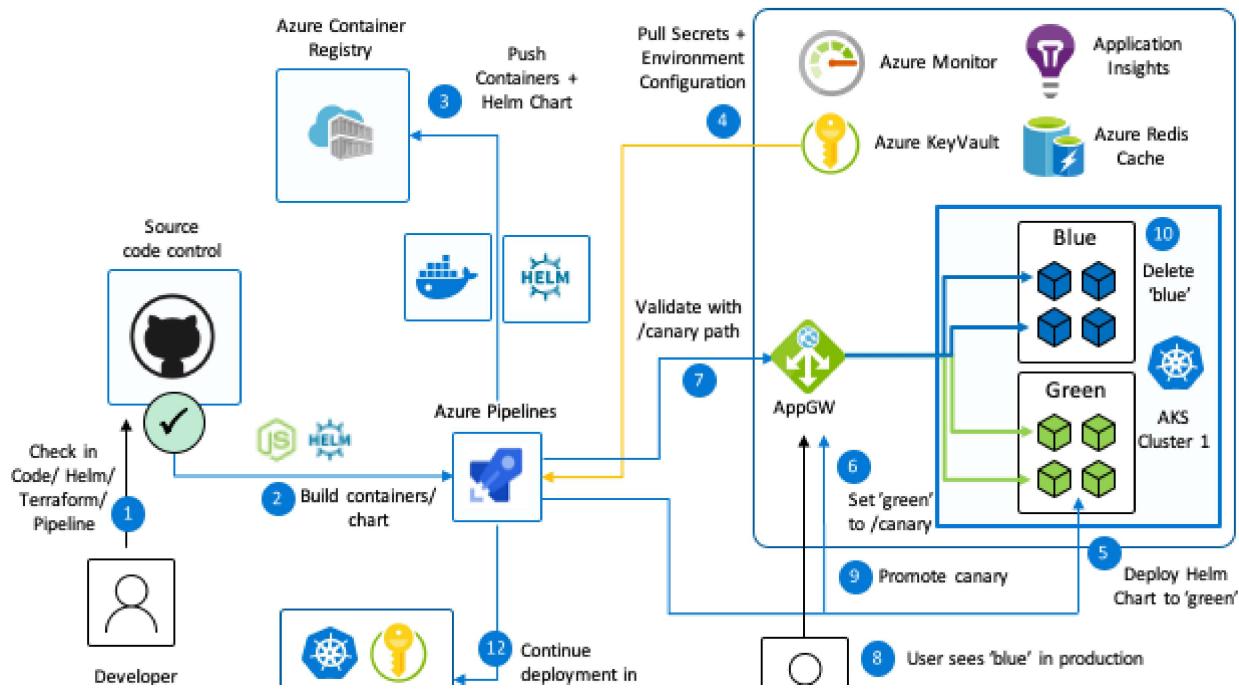
After upgrading the green slot it has now become the production slot for the next iteration.

Since we are using Nginx and Application Insights in our app that means we are getting very detailed metrics, logs and dashboards which will allow us to compare the performance and functionality of our new apps. With a little tuning we can easily set up a Grafana dashboard that allows us to compare both deployments side by side.

That is it for using Nginx!

The next type of deployment will leverage the [azure application gateway as ingress controller](#) to achieve the same. As of today the AKS managed appgateway addon is officially supported but still cannot be automatically deployed via terraform. That means you have to manually activate the addon after you have deployed the cluster as described [here](#).

The process will look like this and is implemented in the [appgw azure pipeline template](#):

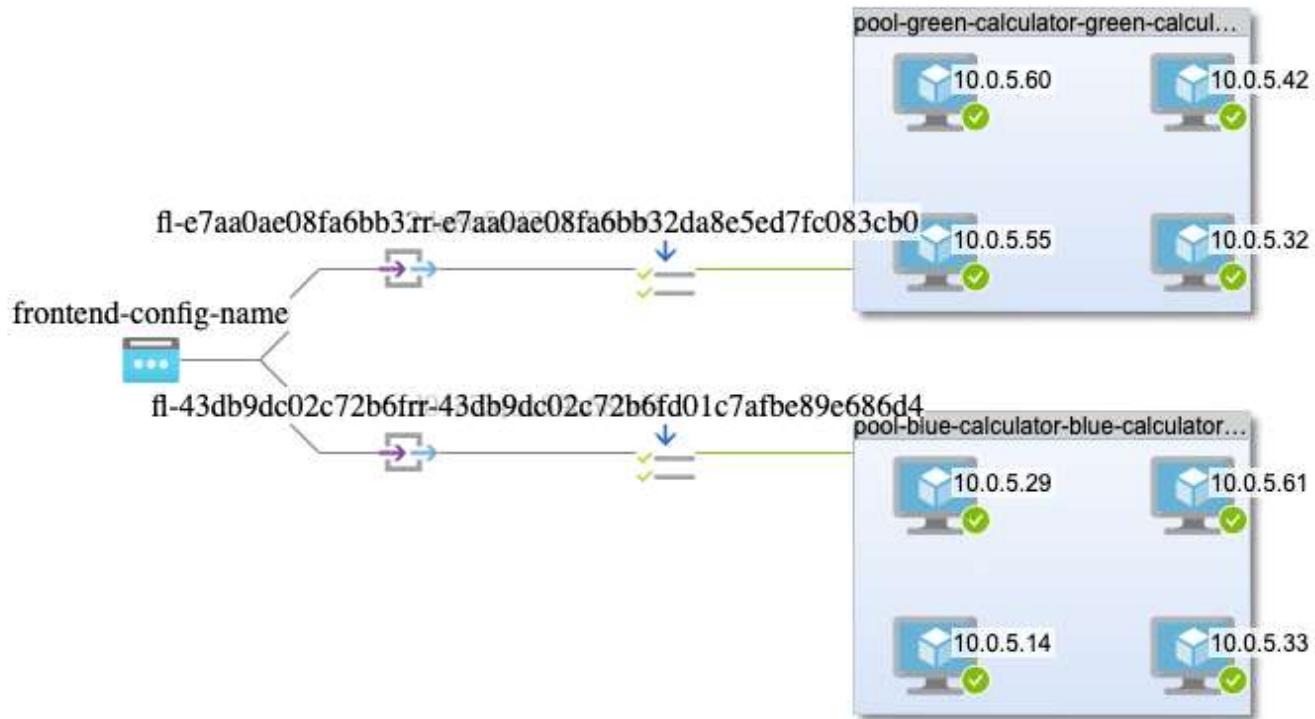




[Open in app](#)

The process is similar to the Nginx deployment, only here we have to use different routes for each deployment.

In comparison to the Nginx ingress controller the AppGateway Ingress controller is running outside of the AKS cluster as its own dedicated managed service, but just as you would expect it will route traffic directly to the individual containers, which is why our application does not have to change to make it work.



Screenshot from the network diagram of the application gateway, where we can see the different backend pools.

The essential process will be the same and controlled by the helm deployment parameters, which will differ only in the ingress.class.

```
scripts > deploy_multicalculator_blue_green_appgw.sh
82 if [ "$GREEN_VERSION" -gt "$BLUE_VERSION" ]; then
83     echo "Green is higher than blue - deploying blue"
84     SLOT="blue"
85 else
86     echo "Blue is higher than green - deploying green"
87     SLOT="green"
88 fi
89
90 if [ "$GREEN_VERSION" == "0" ]; then
91 if [ "$BLUE_VERSION" == "0" ]; then
92     NOCANARY="true"
93 fi
94 fi
95
96 DEPLOY_NAMESPACE=$SLOT-$KUBERNETES_NAMESPACE
```


[Open in app](#)

```
--set image.frontendTag=$BUILD_BUILDDNUMBER --set image.backendTag=$BUILD_BUILDDNUMBER --set image.repository=$AZURE_CONTAINER_REGISTRY_URL
--set dependencies.useAppInsights=true --set dependencies.appInsightsSecretValue=$APPINSIGHTS_KEY --set dependencies.useAzureRedis=true
--set dependencies.redisHostValue=$REDIS_HOST --set dependencies.redisKeyValue=$REDIS_AUTH --set slot=$SLOT --set ingress.class=azure/
application-gateway --set ingress.host=$APPGW_FQDN --wait --timeout 60s
102 else
103 helm upgrade $RELEASE $AZURE_CONTAINER_REGISTRY_NAME/multicalculatorcanary --namespace $DEPLOY_NAMESPACE --install --set replicaCount=2
--set image.frontendTag=$BUILD_BUILDDNUMBER --set image.backendTag=$BUILD_BUILDDNUMBER --set image.repository=$AZURE_CONTAINER_REGISTRY_URL
--set dependencies.useAppInsights=true --set dependencies.appInsightsSecretValue=$APPINSIGHTS_KEY --set dependencies.useAzureRedis=true
--set dependencies.redisHostValue=$REDIS_HOST --set dependencies.redisKeyValue=$REDIS_AUTH --set slot=$SLOT --set ingress.class=azure/
application-gateway --set ingress.host=$APPGW_FQDN --set canary=true --wait --timeout 60s You, a day ago + added tfm
104 fi
```

Same as before we will use ingress objects to register the routing to the right version.

Since the Application Gateway does not support canary routing by HTTP header we are in this case exposing the canary under a different HTTP route named /canary. This allows us same as before validate the canary deployment after it has been deployed and perform a switch between the / — route and the /canary -route after we have confirmed that it works as expected.

charts > multicalculatorcanary > templates > ingress.yaml > [e] {map}

```
35     paths:
36     {{- if and (.Values.canary) (eq .Values.ingress.class "azure/application-gateway") --}}
37       - path: /canary/*
38         backend:
39           serviceName: {{ include "multicalculatorcanary.fullname" . }}-frontend-svc
40           servicePort: {{ .Values.service.port }}
41     {{- else --}}
42       - path:
43         backend:
44           serviceName: {{ include "multicalculatorcanary.fullname" . }}-frontend-svc
45           servicePort: {{ .Values.service.port }}
46     {{- end --}}
47   {{- end }}
```

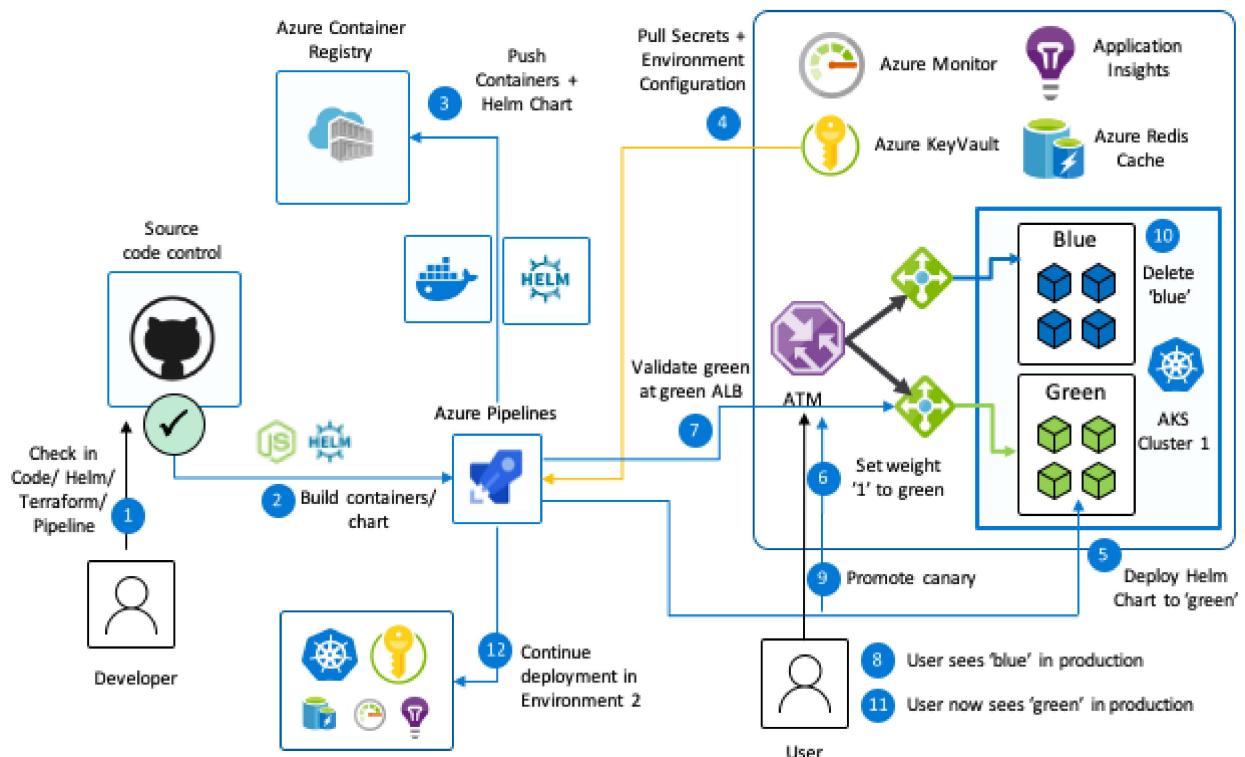
We are switching on the ingress class to publish our app under a different sub path for a canary route

In terms of metrics I have not found a way to compare the traffic monitoring statistics between our two deployment slots, which is a disadvantage, since you have to rely solely on the application insights metrics. However the nice thing about this is that the application gateway can also act as a router between multiple clusters — if they are part of the same VNET. Be aware that the managed Application Gateway Ingress Controller Addon does not work in shared mode, that means if you want one Application Gateway for multiple cluster you have to manually deploy and configure the ingress controller via helm. This makes it possible to perform blue/green deployments not only with our own applications running in different namespaces, but also with blue and green versions running in different clusters. If you are interested in

[Open in app](#)

As final variation we want to take a look at blue/green deployments using Azure Traffic Manager, which is again an azure managed service that works based on DNS resolution. It has the advantage that it will also work on multiple clusters, which do not need to be part of the same VNET or even the same Azure region. However since it is DNS based it will also require is to bring up dedicated Load Balancer IPs for each deployment slot with dedicated DNS entries.

The overall process will look like this:



When using azure traffic manager for blue/green our deployments need to have dedicated DNS entries.

Our instance of azure traffic manager will be configured with the weighted traffic-routing method- which will defines how a HTTP request for our Traffic Manager DNS will get resolved to one of our AKS Load balancer DNS names. This will allow us to decide if the blue or the green version should be preferred by the traffic manager profile.


[Open in app](#)

the endpoints and the routing weight distribution. Our existing shell scripts already have the right spots for these steps and during our terraform deployment we have already set up dedicated load balancer ips and dns entries for both blue and green deployment slots.

```
scripts > deploy_multicalculator_blue_green_tm.sh
105
106 WEIGHT=1
107 if [ "$NOCANARY" == "true" ]; then
108     .. WEIGHT=1000
109     .. helm upgrade $RELEASE $AZURE_CONTAINER_REGISTRY_NAME/multicalculatorcanary --namespace $DEPLOY_NAMESPACE --install --set replicaCount=4
110     .. --set image.frontendTag=$BUILD_BUILDNUMBER --set image.backendTag=$BUILD_BUILDNUMBER --set image.repository=$AZURE_CONTAINER_REGISTRY_URL --set dependencies.useAppInsights=true --set dependencies.appInsightsSecretValue=$APPINSIGHTS_KEY --set dependencies.useAzureRedis=true --set dependencies.redisHostValue=$REDIS_HOST --set dependencies.redisKeyValue=$REDIS_AUTH --set slot=$SLOT --set ingress.enabled=false --set service.type=LoadBalancer --set service.dns=$DNS_LABEL --set service.ip=$IP --set canary=false --wait --timeout 60s
111 else
112     .. helm upgrade $RELEASE $AZURE_CONTAINER_REGISTRY_NAME/multicalculatorcanary --namespace $DEPLOY_NAMESPACE --install --set replicaCount=4
113     .. --set image.frontendTag=$BUILD_BUILDNUMBER --set image.backendTag=$BUILD_BUILDNUMBER --set image.repository=$AZURE_CONTAINER_REGISTRY_URL --set dependencies.useAppInsights=true --set dependencies.appInsightsSecretValue=$APPINSIGHTS_KEY --set dependencies.useAzureRedis=true --set dependencies.redisHostValue=$REDIS_HOST --set dependencies.redisKeyValue=$REDIS_AUTH --set slot=$SLOT --set ingress.enabled=false --set service.type=LoadBalancer --set service.dns=$DNS_LABEL --set service.ip=$IP --set canary=true --wait --timeout 60s
114 fi
115
116 EXISTS=$(az network traffic-manager endpoint show --name $SLOT -g $AKS_GROUP --profile-name $TFM_NAME --type azureEndpoints --query name --output tsv)
117 if [ "$EXISTS" == "$SLOT" ]; then
118     .. az network traffic-manager endpoint update -g $AKS_GROUP --profile-name $TFM_NAME \
119     .. --n $SLOT --type azureEndpoints --target-resource-id $IP_RESOURCE_ID --endpoint-status enabled \
120     .. --weight $WEIGHT --custom-headers host=$DNS
121 else
122     .. az network traffic-manager endpoint create -g $AKS_GROUP --profile-name $TFM_NAME \
123     .. --n $SLOT --type azureEndpoints --target-resource-id $IP_RESOURCE_ID --endpoint-status enabled \
124     .. --weight $WEIGHT --custom-headers host=$DNS
125 fi
```

In our deployment steps we will configure the weight of the endpoints and connect them to our traffic manager profile.

The azure traffic manager will also support more than two public endpoints and allow us to register multiple applications, clusters and instances that can be connected behind the same profile.

[Home](#) > [Resource groups](#) > dzphix_520 >

Name	Status	Monitor status	Type	Weight
green	Enabled	Online	Azure endpoint	1000
blue	Enabled	Online	Azure endpoint	1

The weight and the availability status of each endpoint will decide if traffic will be routed to the individual endpoints.

In terms of metrics the azure traffic manager brings even less metrics than the


[Open in app](#)

slots.

For final comparison I wanted to summarise some of the key aspects of the different implementations and encourage you to try all of them using the samples provided.

	Nginx	Azure AppGW	Azure TrafficManager
Blue/ Green deployments	Yes	Yes	Yes
Fully managed Service	No, managed by you	Yes, as AKS addon for V2	Yes, configured by you
Canary deployment	Paths, Header, Cookies	Paths	No
Connection draining	Yes	Yes	No
Traffic Telemetry	Yes, Grafana	Some, Azure Monitor	No
Works on private IPs	Yes	Yes on V1, not fully private on V2	No
Targeting multiple cluster	No	Yes, but same VNET	Yes, multiple regions

Short summary on the key differences between the technologies used here

I hope you learned something today and I am curious to hear your feedback on how you have implemented your blue/green deployment practice, if my story and samples have helped you to do it and if there is something I have missed.

The end ;)

[Open in app](#)