

Introduction to Crossplane



Pavan Kumar

Feb 27 · 6 min read

How to create any resource on the cloud using Kubernetes manifests and Crossplane.

In the Kubernetes era, all of your application blueprints are packaged into a lot of Kubernetes manifests files or maybe also packages as charts using tools like helm. So how do you create any cloud resource on the cloud? You can maybe use

1. An external terraform module to create the resource.
2. Use a Kubernetes Job and create the resources using AWS SDK's.
3. Use a bash / Python script and internally call AWS CLI commands.

But how reliable is this? Unlike Kubernetes manifests in which the yaml file can be edited on the fly, every time an attribute changes you will have to explicitly call these dependent resources. And in the modern GitOps era having such external dependencies might not be a feasible option for your GitOps Solutions. How do we fix this then? Here comes **Crossplane** into the picture. Crossplane enables you to provision, compose, and consume infrastructure in any **cloud service provider using the Kubernetes API**. Using Crossplane you can create resources on the cloud using simple manifests and can then integrate this with your CI/CD or GitOps pipelines. Crossplane is an open-source project. It is started by Upbound and then later got adopted by the CNCF as a sandbox project.



Image Credits: [Crossplane](#)

What is the entire story all about? (TLDR)

1. Install Cross plane on our Kubernetes Cluster (AKS, GKE, EKS, KIND)
2. Configure Crossplane to communicate with AWS.
3. Install required packages for CrossPlane to communicate with AWS.
4. Create a VPC, SG, RDS using CrossPlane from our Kubernetes Cluster.
5. Verify that the resources have been created from the AWS Console.

Prerequisites

1. A Kubernetes cluster (Can be either On-Prem, AKS, EKS, GKE, Kind).
2. An AWS account.

Story Resources

1. **GitHub Link:** <https://github.com/pavan-kumar-99/medium-manifests>
2. **GitHub Branch:** crossplane

Install Cross plane in a Kubernetes Cluster

You can use an existing Kubernetes cluster for this demo. Alternatively, you can also install a Kubernetes cluster using kind or using GitHub actions. You can refer to my previous articles on how to create a Kubernetes cluster using

1. [GitHub actions](#)
2. [Kind](#)

Resources

1. GitHub Link: <https://github.com/pavan-kumar-99/medium-manifests/tree/crossplane>

Once you have the Kubernetes cluster created let us now install Crossplane in our cluster. You can clone my repo with the crossplane branch for all the manifests used in this article.

###Clone the repo

```
git clone https://github.com/pavan-kumar-99/medium-manifests.git -b crossplane
```

```
cd medium-manifests/crossplane-aws
```

#Create the namespace and install the components using helm

```
kubectl create namespace crossplane-system
```

```
helm repo add crossplane-stable https://charts.crossplane.io/stable  
helm repo update
```

```
helm install crossplane --namespace crossplane-system crossplane-  
stable/crossplane
```

#Check the components are up and healthy

```
kubectl get all -n crossplane-system
```

(OR)

```
git clone https://github.com/pavan-kumar-99/medium-manifests.git -b crossplane
```

```
cd medium-manifests/crossplane-aws
```

```
make install_crossplane
```

Alternatively, you can also use a makefile that I have written. This will install kind in your MAC / Linux machines, create a Kind cluster and then install crossplane in the Kind cluster.

```
1  .PHONY : all install_kind_linux install_kind_mac create_kind_cluster  
2  
3  KIND_VERSION := $(shell kind --version 2>/dev/null)  
4  
5  install_kind_linux :  
6  ifdef KIND_VERSION  
7      @echo "Found version $(KIND_VERSION)"  
8  else  
9      @curl -Lo ./kind https://kind.sigs.k8s.io/dl/v0.10.0/kind-linux-amd64  
10     @chmod +x ./kind  
11     @mv ./kind /bin/kind  
12 endif  
13  
14 install_kind_mac :  
15 ifdef KIND_VERSION  
16     @echo "Found version $(KIND_VERSION)"  
17 else
```

```

18     @brew install kind
19 endif
20
21 create_kind_cluster :
22     @kind create cluster --name crossplane-cluster
23     @kind get kubeconfig --name crossplane-cluster
24     @kubectl config set-context kind-crossplane-cluster
25
26 install_crossplane :
27     @kubectl create namespace crossplane-system
28     @helm repo add crossplane-stable https://charts.crossplane.io/stable
29     @helm repo update
30     @helm install crossplane --namespace crossplane-system crossplane-stable/crossplane
31
32 all : install_kind_linux create_kind_cluster create_kind_cluster

```

Makefile hosted with ❤ by GitHub

[view raw](#)

Let us now install the AWS Provider. This will Install all the CRD's (Custom Resources Definitions) required to create resources on the cloud. Ex:

rdsinstances.database.aws.crossplane.io, ec2.aws.crossplane.io/v1alpha1, etc.

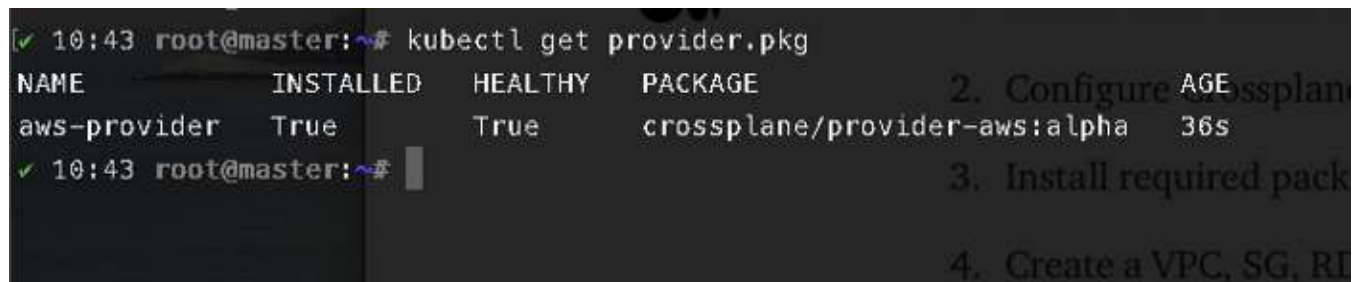
```
kubectl apply -f aws-provider.yaml
```

```

1  apiVersion: pkg.crossplane.io/v1
2  kind: Provider
3  metadata:
4    name: aws-provider
5  spec:
6    package: crossplane/provider-aws:alpha

```

aws-provider.yaml hosted with ❤ by GitHub

[view raw](#)


```

✓ 10:43 root@master:~# kubectl get provider.pkg
NAME          INSTALLED  HEALTHY  PACKAGE                                AGE
aws-provider  True       True     crossplane/provider-aws:alpha         36s
✓ 10:43 root@master:~#

```

Provider Package

###Once you install the Provider, wait for the Provider to be healthy by executing

```
kubectl get provider.pkg
```

Once the Provider is healthy let us now configure the Provider to communicate with AWS by creating a `ProviderConfig` definition. Make sure that you have already configured your credentials using `AWS configure` (From the cli, if you are running the commands from a local cluster).

###Generate the configuration files with the AWS Credentials.

```
AWS_PROFILE=default && echo -e "[default]\naws_access_key_id = $(aws\nconfigure get aws_access_key_id --profile\n$AWS_PROFILE)\naws_secret_access_key = $(aws configure get\naws_secret_access_key --profile $AWS_PROFILE)" > creds.conf
```

###Create a Kubernetes secret with the configuration file generated.

```
kubectl create secret generic aws-secret-creds -n crossplane-system\n--from-file=creds=./creds.conf
```

###Once the secret is created let us now create the Provider config for our AWS account.

```
kubectl apply -f provider-config.yaml
```

```
1  apiVersion: aws.crossplane.io/v1beta1
2  kind: ProviderConfig
3  metadata:
4    name: awsconfig
5  spec:
6    credentials:
7      source: Secret
8      secretRef:
9        namespace: crossplane-system
10       name: aws-secret-creds
11       key: creds
```

provider-config.yaml hosted with ❤ by GitHub

[view raw](#)

Upon successful creation, your local cluster should now be able to communicate with AWS. Let us now try creating the following scenario. Let us create a VPC and a Security

Group that would allow access from Port 3306 from anywhere from the world. Let us simultaneously create an RDS and attach the aforementioned SG to the same RDS Instance so that it would be publically accessible. Once this resource is created we will create a pod in our local cluster and check if it can access the RDS Instance. Seems good? Let us now get into action.

Let us create a VPC in the us-east-1 region with the below-mentioned spec.

```
1  apiVersion: ec2.aws.crossplane.io/v1beta1
2  kind: VPC
3  metadata:
4    name: production-vpc
5  spec:
6    forProvider:
7      region: us-east-1
8      cidrBlock: 192.168.0.0/16
9      enableDnsSupport: true
10     enableDnsHostNames: true
11     tags:
12     - key: Environment
13       value: Production
14     - key: Owner
15       value: Pavan
16     - key: Name
17       value: production-vpc
18     instanceTenancy: default
19   providerConfigRef:
20     name: awsconfig
```

aws-vpc.yaml hosted with ❤ by GitHub

[view raw](#)

```
kubectl apply -f aws-vpc.yaml
```

###Let us check the status of the VPC. We are now referring to the provider created earlier in (line no 20).

```
kubectl get vpc
```

```
✓ 16:29 root@master:~# k get vpc
NAME                READY   SYNCED   ID                               CIDR           AGE
production-vpc      True    True     vpc-03bd0495614a9ab67          192.168.0.0/16 46s
✓ 16:29 root@master:~#
```

AWS VPC for CrossPlane

```
apiVersion: ec2.aws.crossplane.io/v1beta1
kind: VPC
metadata:
```

VPC created with VPC ID and CIDR block and the sync and Ready state

Once our VPC is successfully created let us create 2 subnets and attach an internet gateway to our VPC and also add a Route table for the same so that we can create our RDS in these Public subnets and then access from our local pod. However this is not the suggested method in Production, you should never spin your RDS in a Public Subnet in a Production environment.

```
1  apiVersion: ec2.aws.crossplane.io/v1beta1
2  kind: Subnet
3  metadata:
4    name: prod-subnet-1
5  spec:
6    forProvider:
7      region: us-east-1
8      availabilityZone: us-east-1a
9      cidrBlock: 192.168.1.0/24
10     vpcIdRef:
11       name: production-vpc
12     tags:
13       - key: Environment
14         value: Production
15       - key: Name
16         value: prod-subnet-1
17       - key: Owner
18         value: Pavan
19     mapPublicIPOnLaunch: true
20   providerConfigRef:
21     name: awsconfig
22
23 ---
24
25 apiVersion: ec2.aws.crossplane.io/v1beta1
26 kind: Subnet
27 metadata:
28   name: prod-subnet-2
29 spec:
30   forProvider:
31     region: us-east-1
32     availabilityZone: us-east-1b
33     cidrBlock: 192.168.2.0/24
34     vpcIdRef:
35       name: production-vpc
36   tags:
37     - key: Environment
38       value: Production
```

```

38     value: production
39   - key: Name
40     value: prod-subnet-2
41   - key: Owner
42     value: Pavan
43   mapPublicIPOnLaunch: true
44   providerConfigRef:
45     name: awsconfig

```

aws-subnet.yaml hosted with  by GitHub

[view raw](#)

```
kubectl apply -f aws-subnet.yaml
```

###Let us check the status of the subnets.

```
kubectl get subnets
```

Let us now create the corresponding Internet gateway and Route table.

```

1  apiVersion: ec2.aws.crossplane.io/v1beta1
2  kind: InternetGateway
3  metadata:
4    name: production-internetgateway
5  spec:
6    forProvider:
7      region: us-east-1
8      vpcIdRef:
9        name: production-vpc
10   providerConfigRef:
11     name: awsconfig
12
13   ---
14
15
16  apiVersion: ec2.aws.crossplane.io/v1alpha4
17  kind: RouteTable
18  metadata:
19    name: production-routetable
20  spec:
21    forProvider:
22      region: us-east-1
23      routes:
24        - destinationCidrBlock: 0.0.0.0/0
25          gatewayIdRef:
26            name: production-internetgateway
27      associations:

```



```

28     - subnetIdRef:
29       name: prod-subnet-2
30     - subnetIdRef:
31       name: prod-subnet-1
32     vpcIdRef:
33       name: production-vpc
34   providerConfigRef:
35     name: awsconfig

```

aws-igwrt.yaml hosted with ❤ by GitHub

[view raw](#)

```
kubectl apply -f aws-igwrt.yaml
```

###Let us check the status of the Route table and Internet Gateway

```
kubectl get InternetGateway,RouteTable
```

Let us now create the security group that would allow communication over port 3306 to the Internet. Later we will attach this security group to our RDS Instance.

```

1  apiVersion: ec2.aws.crossplane.io/v1beta1
2  kind: SecurityGroup
3  metadata:
4    name: rds-access-sg
5  spec:
6    forProvider:
7      region: us-east-1
8      vpcIdRef:
9        name: production-vpc
10     groupName: mysql-sg
11     description: RDS communication to local application Pods
12     tags:
13     - key: Environment
14       value: Production
15     - key: Owner
16       value: Pavan
17     - key: Name
18       value: rds-access-sg
19     ingress:
20     - fromPort: 3306
21       toPort: 3306
22       ipProtocol: tcp
23       ipRanges:
24       - cidrIp: 0.0.0.0/0
25     providerConfigRef:

```

2b name: awsconfig

aws-sg.yaml hosted with ❤ by GitHub

[view raw](#)

```
kubectl apply -f aws-sg.yaml
```

###Let us check the status of the Route table and Internet Gateway

```
kubectl get SecurityGroup
```

```
16:35 root@master:~# k get securitygroups
NAME          READY   SYNCED   ID              VPC              AGE
rds-access-sg True     True     sg-0441d19ed64bdefcb vpc-03bd0495614a9ab67 28s
```

Security group created in the afore-mentioned VPC.

Let us now create the RDS instance, but before we do that we would need a subnet group in which the RDS instance has to be created. We will use the subnets created earlier in the DB Subnet Group.

```
1  apiVersion: database.aws.crossplane.io/v1beta1
2  kind: DBSubnetGroup
3  metadata:
4    name: prod-subnet-group
5  spec:
6    forProvider:
7      region: us-east-1
8      description: "Prod Subnet group"
9      subnetIdRefs:
10       - name: prod-subnet-2
11       - name: prod-subnet-1
12    providerConfigRef:
13      name: awsconfig
14
15  ---
16
17  apiVersion: database.aws.crossplane.io/v1beta1
18  kind: RDSInstance
19  metadata:
20    name: production-rds
21  spec:
22    forProvider:
23      allocatedStorage: 50
24      autoMinorVersionUpgrade: true
25      applyModificationsImmediately: false
```

```
--
26 backupRetentionPeriod: 0
27 caCertificateIdentifier: rds-ca-2019
28 copyTagsToSnapshot: false
29 dbInstanceClass: db.t2.small
30 dbSubnetGroupName: prod-subnet-group
31 vpcSecurityGroupIDRefs:
32   name: rds-access-sg
33 deletionProtection: false
34 enableIAMDatabaseAuthentication: false
35 enablePerformanceInsights: false
36 engine: mysql
37 region: us-east-1
38 engineVersion: 5.6.35
39 finalDBSnapshotIdentifier: muvaf-test
40 licenseModel: general-public-license
41 masterUsername: admin
42 multiAZ: true
43 port: 3306
44 tags:
45   - key: Name
46     value: production-rds
47   - key: Environment
48     value: Production
49   - key: Owner
50     value: Pavan
51 publiclyAccessible: true
52 storageEncrypted: false
53 storageType: gp2
54 providerConfigRef:
55   name: awsconfig
56 writeConnectionSecretToRef:
57   name: production-rds-conn-string
58   namespace: default
```

```
kubectl apply -f aws-rds.yaml
```

###Let us check the status of the RDS Instance. The credentials are stored in a secret called production-rds-conn-string in the default namespace. (line no 56)

```
kubectl get RDSInstance
```





The AWS Infra is now ready to serve us.....

Now let us try to access our Mysql RDS Instance. We can access this by decoding the secret **production-rds-conn-string** created in the default namespace. You can connect to the database using the MySQL client

```
mysql -h <hostname> -u <user_name> -p <password>
```

Alternatively, you can spin up a pod and connect from the pod itself.

```
1  apiVersion: v1
2  kind: Pod
3  metadata:
4    name: mysql-rds-connection-test
5    namespace: default
6  spec:
7    containers:
8      - name: mysqlconn
9        image: mysql:latest
10       command: ['mysql']
11       args: ['-c', 'show databases;']
12       env:
13         - name: MYSQL_HOST
14           valueFrom:
15             secretKeyRef:
16               name: production-rds-conn-string
17               key: endpoint
18         - name: MYSQL_PWD
19           valueFrom:
20             secretKeyRef:
21               name: production-rds-conn-string
22               key: password
23         - name: USER
24           valueFrom:
25             secretKeyRef:
26               name: production-rds-conn-string
27               key: username
```

```
28   - name: MYSQL_TCP_PORT
29     valueFrom:
30       secretKeyRef:
31         name: production-rds-conn-string
32         key: port
```

aws-rds-connection-test.yaml hosted with ❤ by GitHub

[view raw](#)

```
###Create a testpod that shows all the databases in the RDS Instance
kubectl apply -f aws-rds-connection-test.yaml
```

You should now see the databases in the logs of the Pod.



Throughout the article we have hardcoded the names, we also have an option to filter the resources using tags. But for some reason, my resources were not being filtered even after tagging them. If you were able to filter them using the tags please feel free to paste the solution in the comments section.

Until we meet again.....

Conclusion

Thanks for reading my article. Hope you have liked it. Here are some of my other articles that may interest you.

Recommended

<p>Introduction to Jenkins Operator</p> <p>Getting started with Jenkins Operator in Kubernetes</p> <p>medium.com</p>	
<p>Introduction to Bitnami Sealed Secrets</p> <p>How to store your secrets in GitHub using Sealed Secrets and Kubese</p> <p>medium.com</p>	
<p>Introduction to External DNS in Kubernetes</p> <p>How to automatically create DNS records in Kubernetes using External DNS</p> <p>medium.com</p>	
<p>Creating a GKE Cluster with GitHub Actions</p> <p>Automating Kubernetes Cluster creation and Bootstrapping using GitHub Actions</p> <p>medium.com</p>	

Reference

<p>crossplane/crossplane</p> <p>Crossplane, a Cloud Native Computing Foundation sandbox project, is an open-source Kubernetes add-on that extends any...</p> <p>github.com</p>	
---	--



Buy me a coffee

[Crossplane](#)

[Kubernetes](#)

[AWS](#)

[Cloud Computing](#)

[Github](#)

[About](#) [Help](#) [Legal](#)

Get the Medium app

